

---

# *Redes de Computadores*

## **The Network Layer**

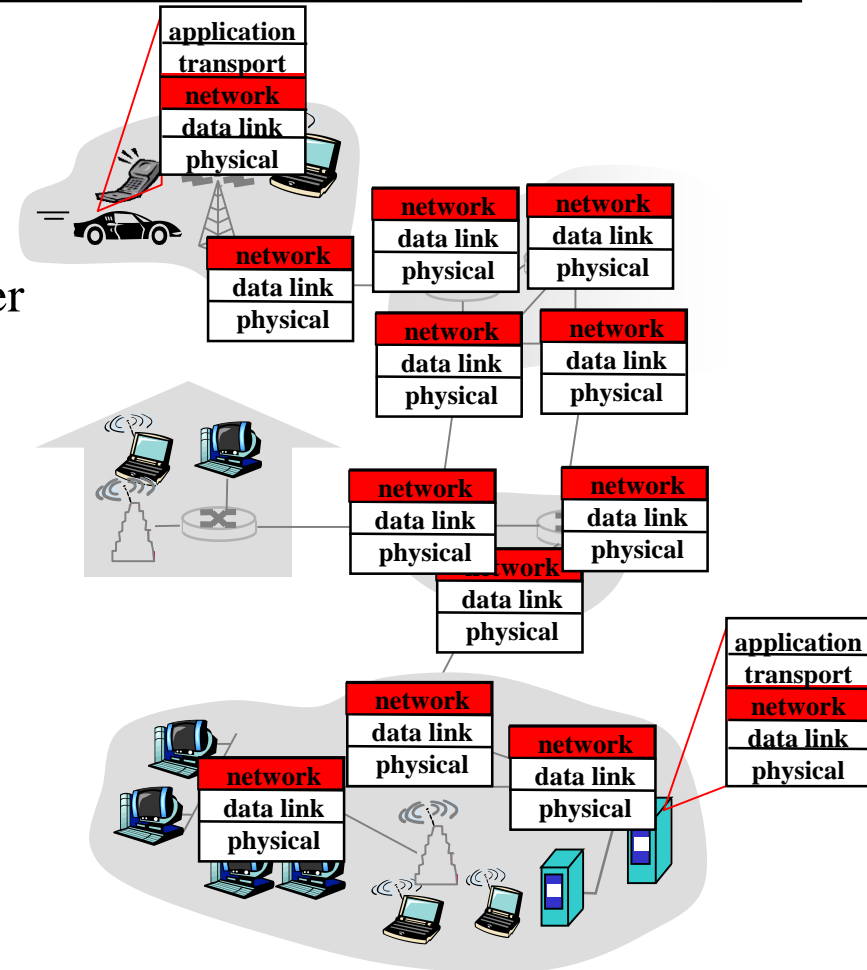
*Manuel P. Ricardo*

*Faculdade de Engenharia da Universidade do Porto*

- 
- » What are the main functions of the network layer?
  - » What are the differences between virtual circuit and datagram networks?
  - » How is forwarding handled in both type of networks?
  - » What are the main functions of a router?
  - » What are the formats of IP addresses?
  - » How to form subnets?
  - » What services are provided by ARP, ICMP, DHCP and NAT? How do these protocols work?
  - » What are differences the between IPv4 and IPv6?

# Network Layer Overview

- ◆ Network layer
  - » transports packets (datagrams)
  - » from sending host to receiving host
  - » functions located in every host and router
- ◆ Sender
  - » encapsulates transport data into packets
  - » generates packets
- ◆ Receiver
  - » receives packets
  - » delivers data to transport layer
- ◆ Router
  - » Receives packets from input line
  - » examines network layer header
  - » forwards packets through adequate output line(s)



# *Network Layer – Main Functions*


---

## ◆ Forwarding

- » router forwards packet from input port to output port

## ◆ Routing

- » determine route taken by packets, from source to destination
- » algorithms, shortest path



---

# *Virtual Circuits and Datagram Networks*


# *Network Layer – Connection and Connectionless Service*

---


## Services provided by network layer

- » Datagram network  
➔ connectionless service
- » Virtual Circuit network  
➔ connection oriented service

(b) Virtual circuit packet switching



(c) Datagram packet switching




# *Virtual Circuit (VC)*

---

- ◆ Phases

circuit establishment → data transference → circuit termination

- ◆ Packet carries identifier of Virtual Circuit



- ◆ Path defined from source to destination


sequence of VC identifiers, one for each link along path

- ◆ Router

- » maintains “state” for every supported circuit
- » may allocate resources (bandwidth, buffers) per Virtual Circuit

# VC - Forwarding Table

VC number



Forwarding table in  
northwest router:

Incoming interface	Incoming VC #	Outgoing interface	Outgoing VC #
1	12	3	22
2	63	1	18
3	7	2	17
1	97	3	87
...	...	...	...


Routers maintain connection state information!



# *Datagram Networks*

---

- ◆ No circuit establishment; no circuit concept
- ◆ Packets
  - » forwarded using destination host address
  - » packets between same source-destination pair may follow different paths



# *Forwarding Table*

---

<u>Destination Address Range</u>	<u>Output Link Interface</u>
11001000 00010111 00010000 00000000 through 11001000 00010111 00010111 11111111	0
11001000 00010111 00011000 00000000 through 11001000 00010111 00011000 11111111	1
11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111	2
otherwise	3

$2^{32}$  possible entries in IPv4

- 
- ◆ How to reduce the number of entries in the forwarding table?

# *Longest Prefix Matching*

---

<u>Prefix Match</u>	<u>Link Interface</u>
11001000 00010111 00010	0
11001000 00010111 00011000	1
11001000 00010111 00011	2
otherwise	3

## Examples. Which Interface?

DA: 11001000 00010111 00010110 10100001 → 0

DA: 11001000 00010111 00011000 10101010 → 1,2 → 1

  
longest prefix

# *Virtual-Circuit versus Datagram Networks*

---

Issue	Datagram subnet	Virtual-circuit subnet
Circuit setup	Not needed	Required
Addressing	Each packet contains the full source and destination address	Each packet contains a short VC number
State information	Routers do not hold state information about connections	Each VC requires router table space per connection
Routing	Each packet is routed independently	Route chosen when VC is set up; all packets follow it
Effect of router failures	None, except for packets lost during the crash	All VCs that passed through the failed router are terminated
Quality of service	Difficult	Easy if enough resources can be allocated in advance for each VC
Congestion control	Difficult	Easy if enough resources can be allocated in advance for each VC

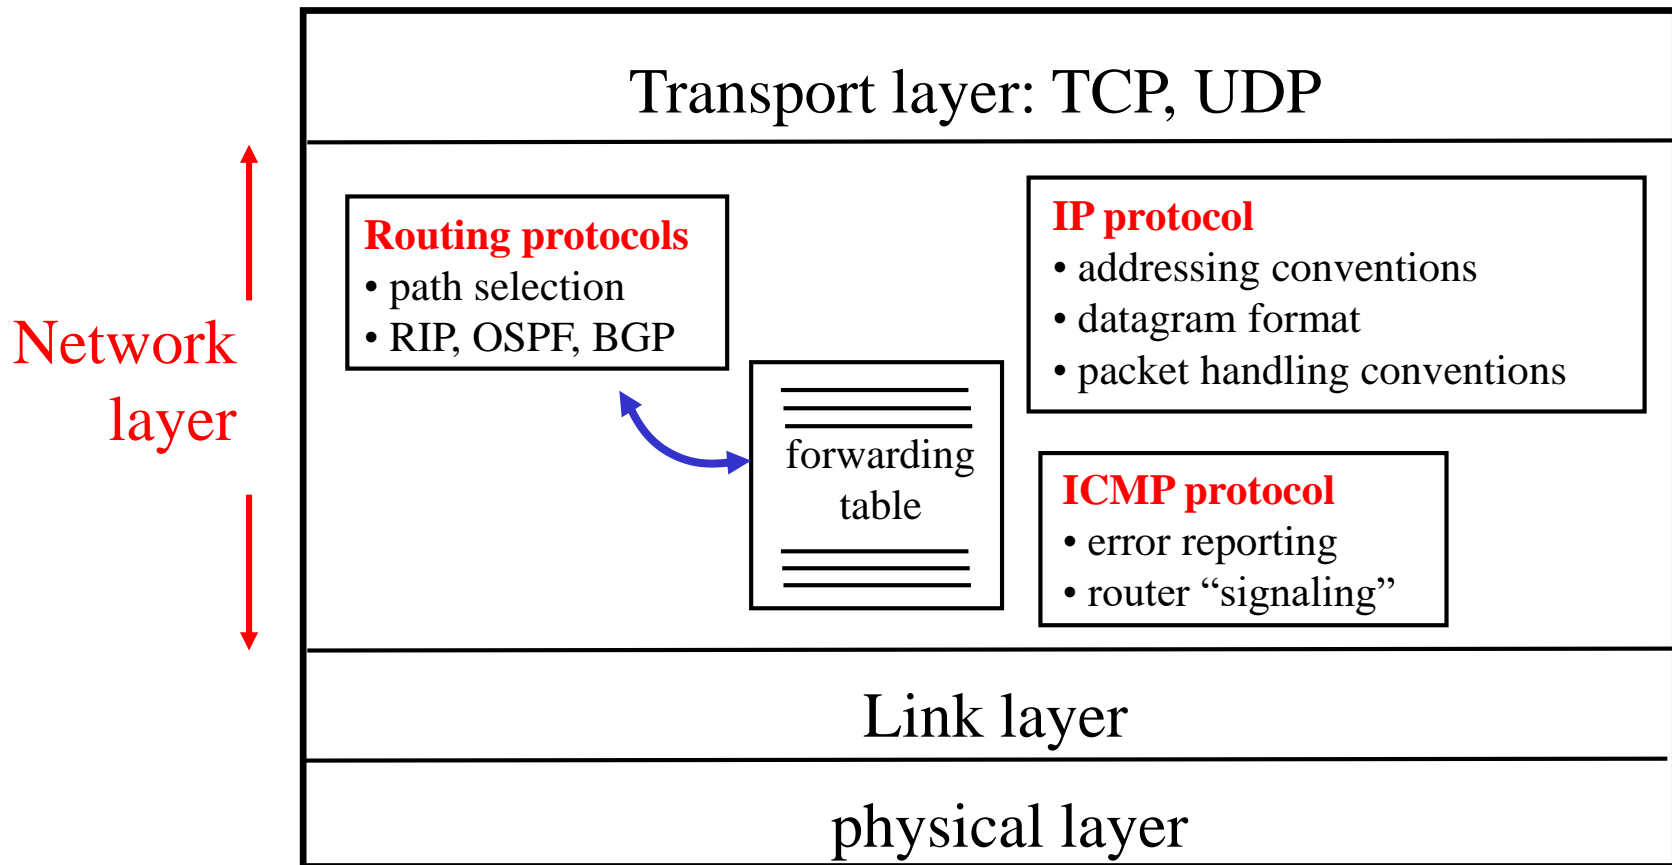
---

# *Internet Protocol*

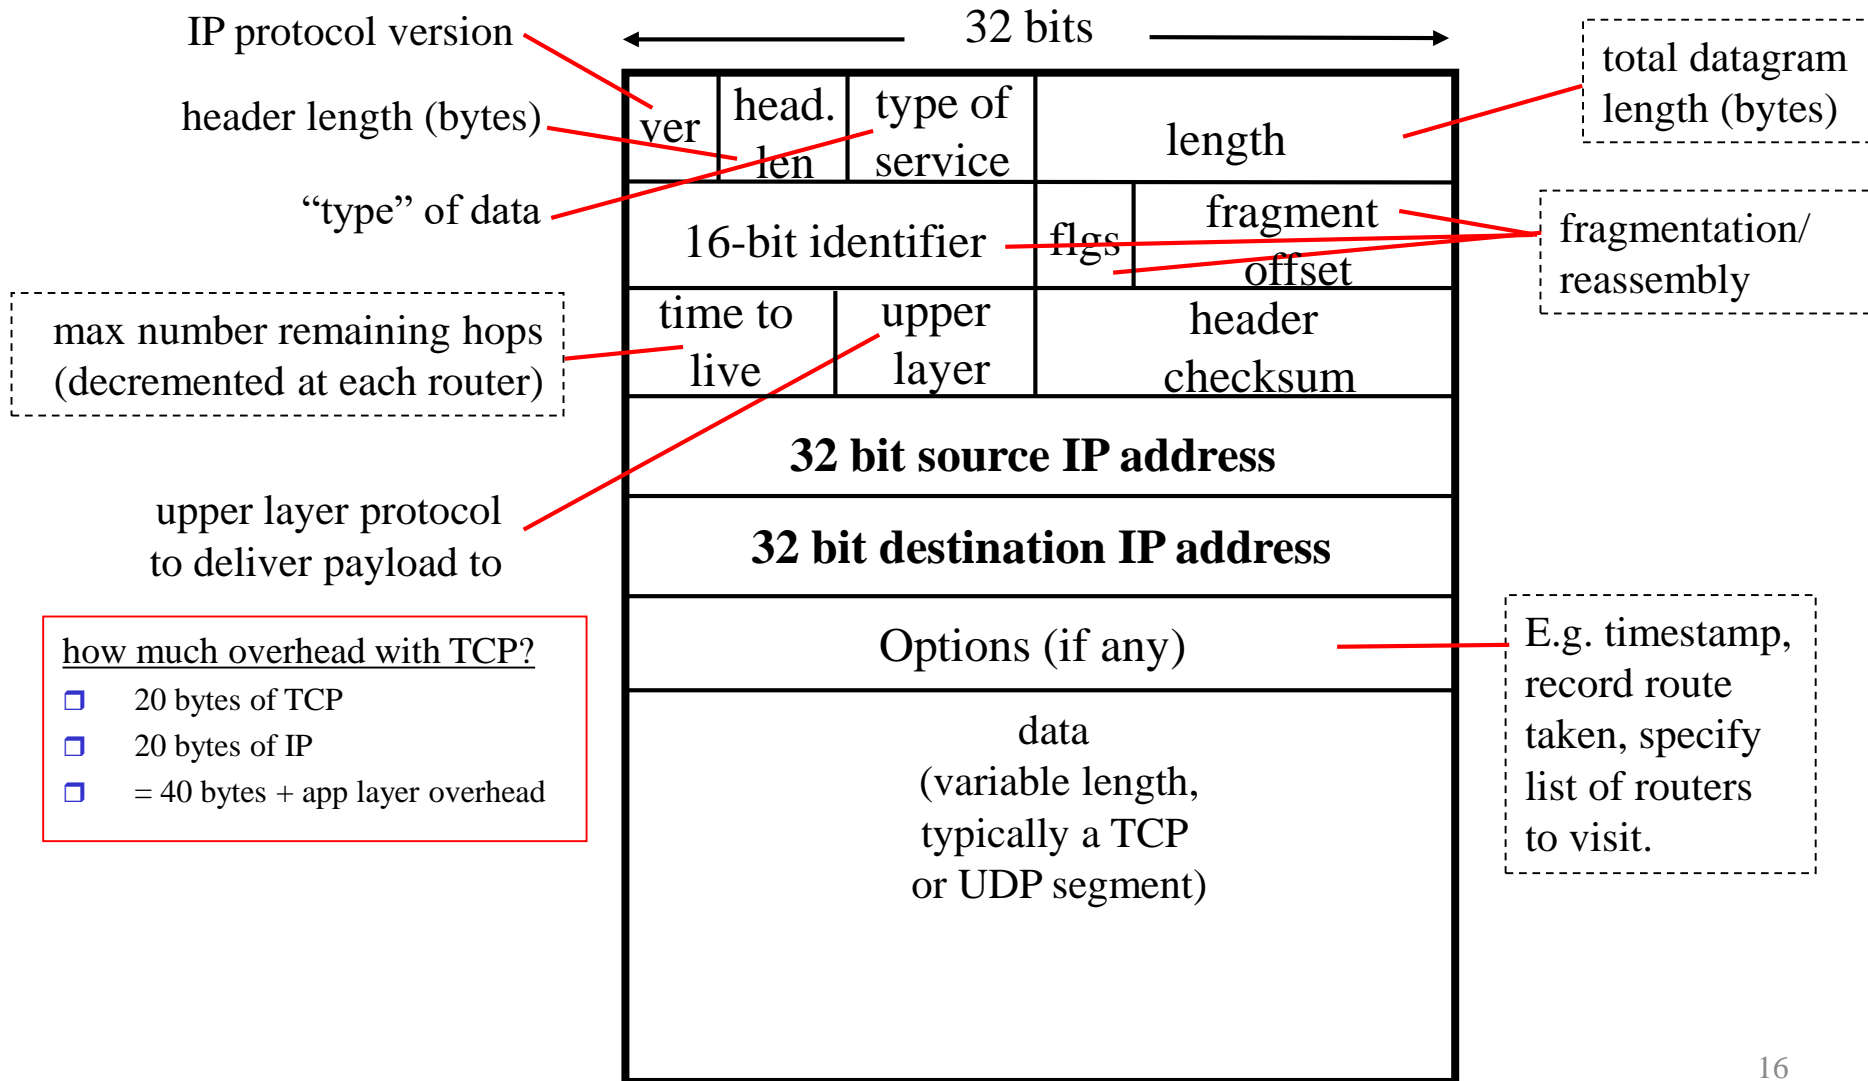
# *The Internet Network layer*

---

Host, router network layer functions



# IP Datagram Format





# Internet Checksum

- ◆ The Internet (**not layer 2**) uses a checksum
  - » easily implementable in software →
  - » 1's complement sum of 16 bit words
  - » Performance: d=2

```
u_short
cksum(u_short *buf, int count)
{
    register u_long sum = 0;


    while (count--)
    {
        sum += *buf++;
        if (sum & 0xFFFF0000)
        {
            /* carry occurred,
             so wrap around */
            sum &= 0xFFFF;
            sum++;
        }
    }
    return ~(sum & 0xFFFF);
}
```

- ◆ One's complement sum
  - » Mod-2 addition **with carry-out**
  - » Carry-out in the most-significant-bit is added to the least-significant bit
  - » Get one's complement of "one's complement sum"

	1010011
	0110110
carry-out ①	0001001
Carry wrap-around	0000001
	0001010
<b>One's complement =</b>	<b>1110101</b>

# *IP Fragmentation and Reassembly*

- ◆ Network links have MTU
  - » MTU - max. transfer size
  - » largest possible link-level frame
  - » different link types, different MTUs
- ◆ Large IP datagram is fragmented
  - » one datagram → n datagrams
  - » “reassembled” at final destination
  - » IP header bits used to identify, order related fragments



# IP Fragmentation and Reassembly

## Example

- ❑ 4000 byte datagram
- ❑ 3980 bytes data + 20 bytes IP header
- ❑ MTU = 1500 bytes

	length =4000	ID =x	fragflag =0	offset =0	
--	-----------------	----------	----------------	--------------	--

One large datagram becomes several smaller datagrams

1480 bytes in data field

offset =  
 $1480/8$

	length =1500	ID =x	fragflag =1	offset =0	
--	-----------------	----------	----------------	--------------	--

	length =1500	ID =x	fragflag =1	offset =185	
--	-----------------	----------	----------------	----------------	--

	length =1040	ID =x	fragflag =0	offset =370	
--	-----------------	----------	----------------	----------------	--

# *IP Addressing - Introduction*


---

## ◆ IP address

- » 32-bit identifier for host/router interface

## ◆ Interface

- » connection between host/router and physical link
- » Routers have multiple interfaces
- » IP addresses associated with interface



# Subnets

---


- ◆ IP address

- » subnet part → high order bits
- » host part → low order bits

- ◆ Subnet → set of interfaces


- » with same subnet part of IP address
- » can reach each other without router intervention

Network consisting of 3 subnets



# 6 Subnets

---




# *IP Addressing - CIDR*

---

## CIDR: Classless InterDomain Routing

- » subnet portion of address has arbitrary length
- » address format → a.b.c.d/x
  - where x is # bits in subnet portion of address



200.23.16.0/23

# *Special IP Addresses*

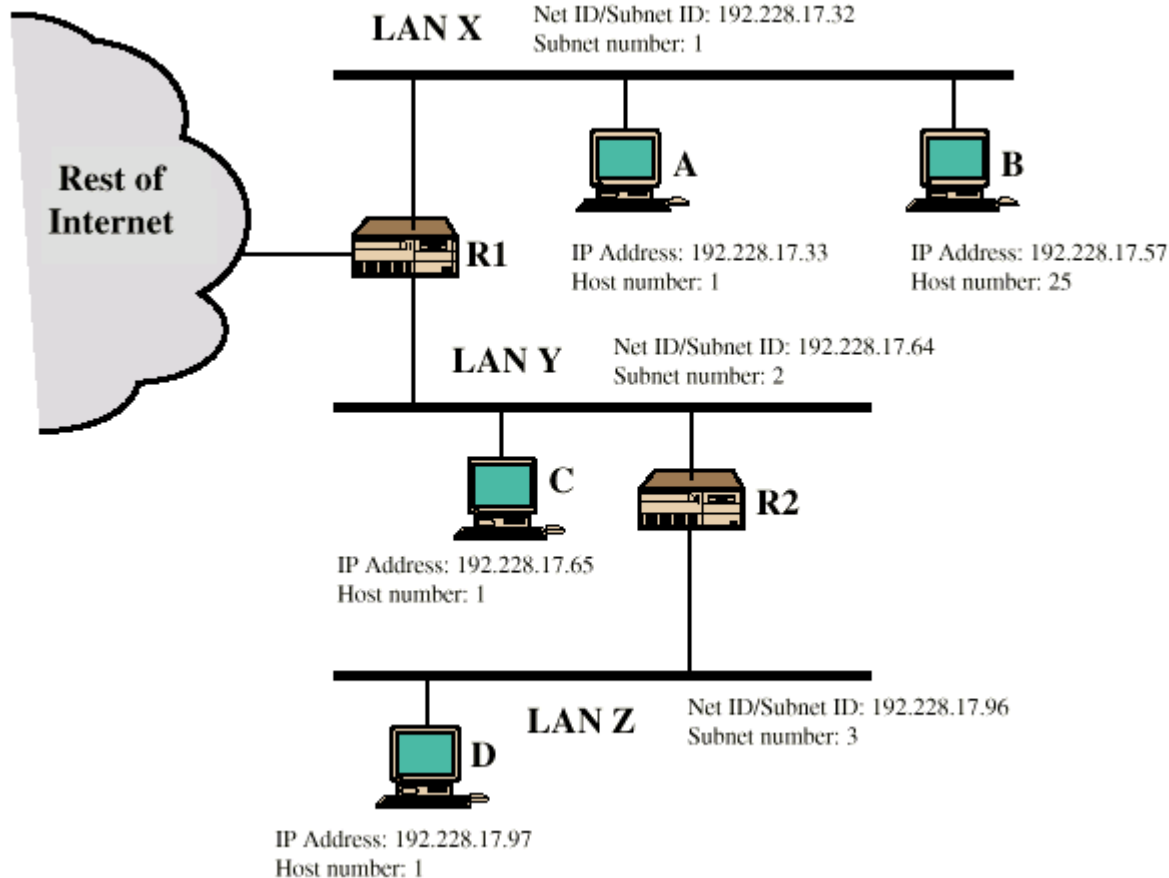
---

0 0																														This host										
0 0										...										0 0										Host	A host on this network									
1 1																														Broadcast on the local network										
Network										1 1 1 1										...										1 1 1 1										Broadcast on a distant network
127					(Anything)																									Loopback										



# Forming Sub-Networks

Network **192.228.17.0/24** is divided in **8 subnetworks** → masks of 27 bits



Subnetwork mask – 27 bits

*subnetid – 3 bits (8 subnetworks)*

11000000 11100100 00010001 **011**00000  
192.228.17.96/27

*hostid – 5 bits*

*30 hosts per subnet supported*  
*all 0 – identifies subnet*  
*all 1 – broadcast address*

## Example of subnetworks

192.228.17.0/27 (.00000000)

192.228.17.32/27 (.00100000)

192.228.17.64/27 (.01000000)

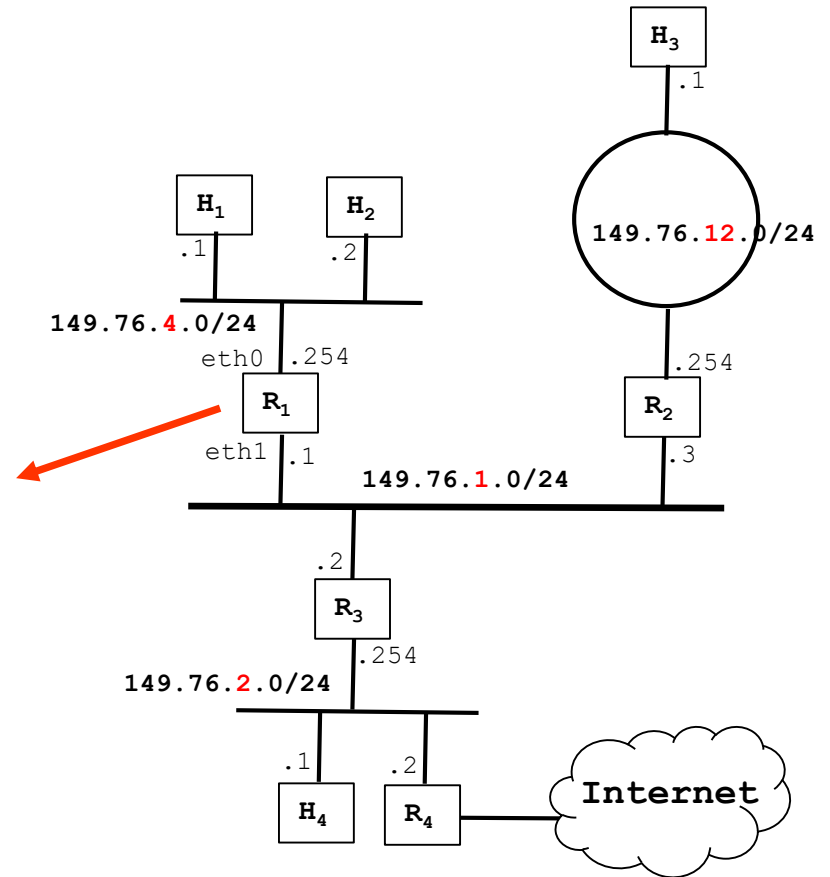
192.228.17.96/27 (.01100000)

....

192.228.17.224/27 (.11100000)

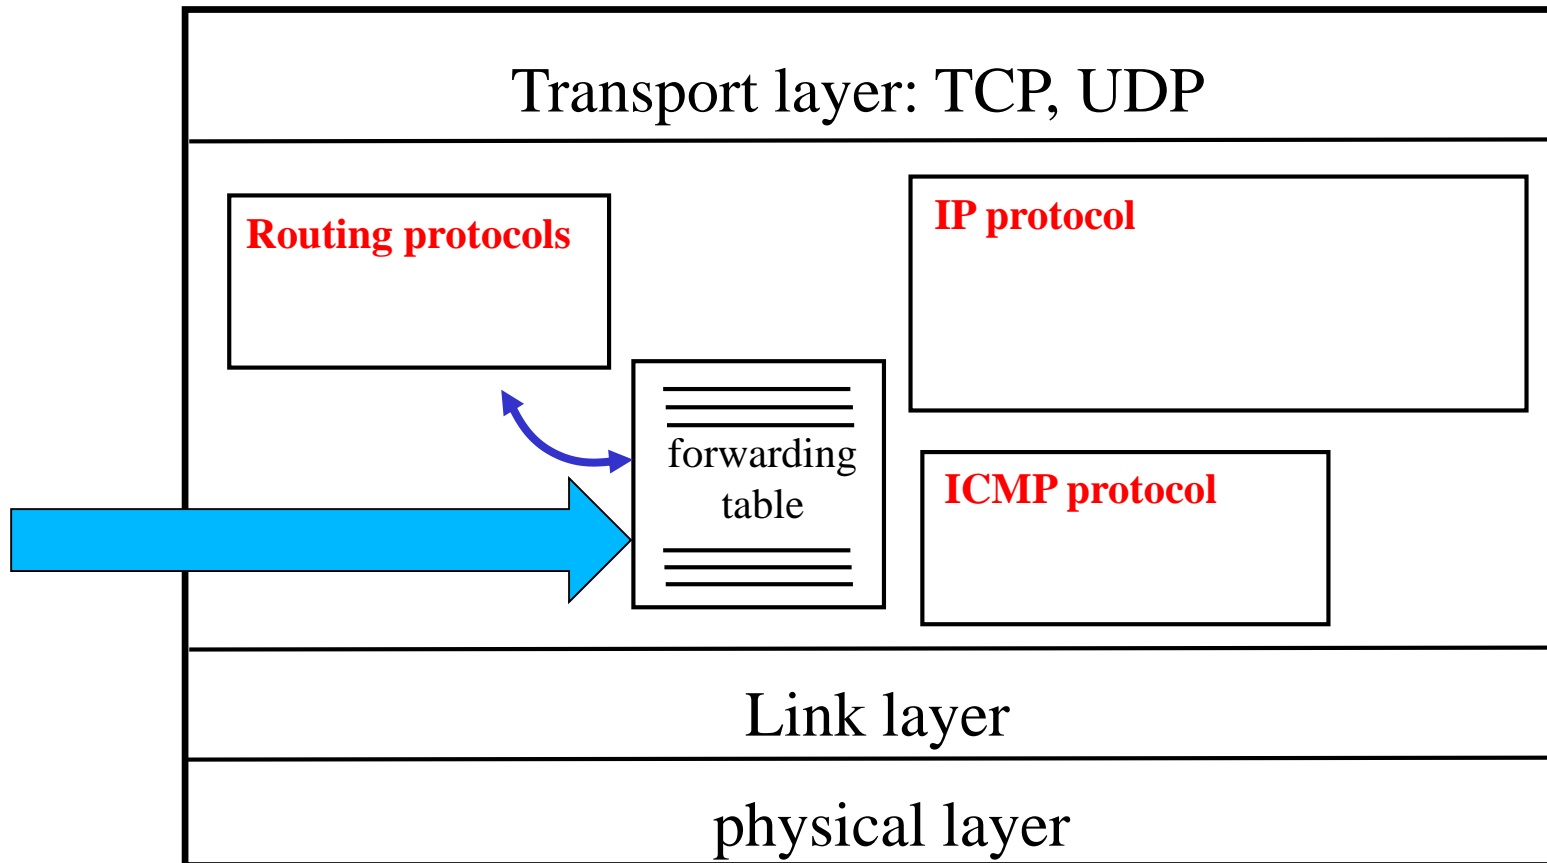
# Forwarding Table at R1

Destination	Gateway	Interface
149.76.1.0/24	-	eth1
149.76.2.0/24	149.76.1.2	eth1
149.76.4.0/24	-	eth0
149.76.12.0/24	149.76.1.3	eth1
0/0	149.76.1.2	eth1



# *Forwarding Table at R1*

---



- 
- ♦ What is a loopback interface? What is its IP address?

# *IP Forwarding Function*

---

- ◆ Forwarding table has entries in format

`<networkAddress/mask, port>`

- ◆ Forwarding function

- » When a datagram arrives with destination address **A**, then

- For each entry of the forwarding table
  - `val = A & mask*` (e.g., `mask=8, mask*=255.0.0.0`)
  - `if ( val == networkAddress & mask* )`
    - add corresponding output port to the set of candidate ports
- Select the port with the largest mask → most specific route

- » Example

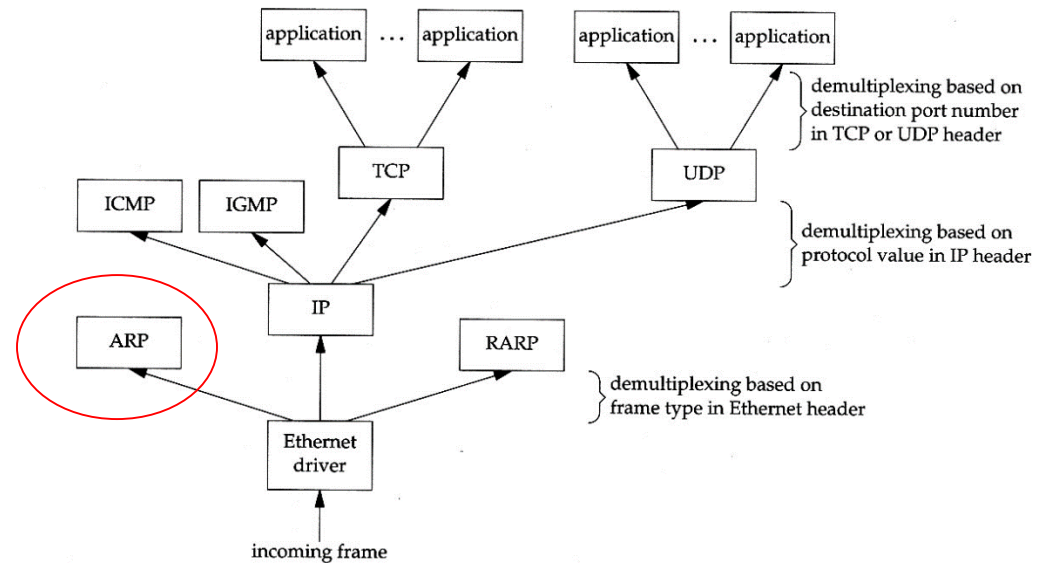
- `frdTbl = { <128.32.1.5/16, 1>, <128.32.225.0/18, 3>, <128.0.0.0/8, 5> }`
- Datagram with destination address `A=128.32.195.1`
- Set of candidate output ports → `{1, 3, 5}`.
- Selected port → **3** ← largest mask, 18 bits

---

# *Address Resolution Protocol*

# Demultiplexing

- ◆ Ethernet header (type)
  - » IP - 0x0800
  - » ARP - 0x0806
  - » RARP - 0x8035
  - » IPX- 0x8037
  - » IPv6 - 0x86DD
  - » MPLS - 0x8847
- ◆ IP header (protocol)
  - » ICMP - 1
  - » IGMP - 2
  - » TCP - 6
  - » UDP - 17
- ◆ TCP/UDP header (port)
  - » FTP - 21
  - » Telnet - 23
  - » HTTP - 80
  - » SMTP - 25



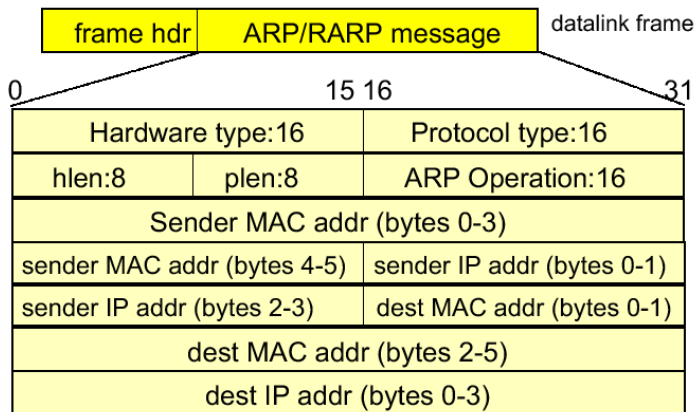
# *ARP – Address Resolution Protocol*

---

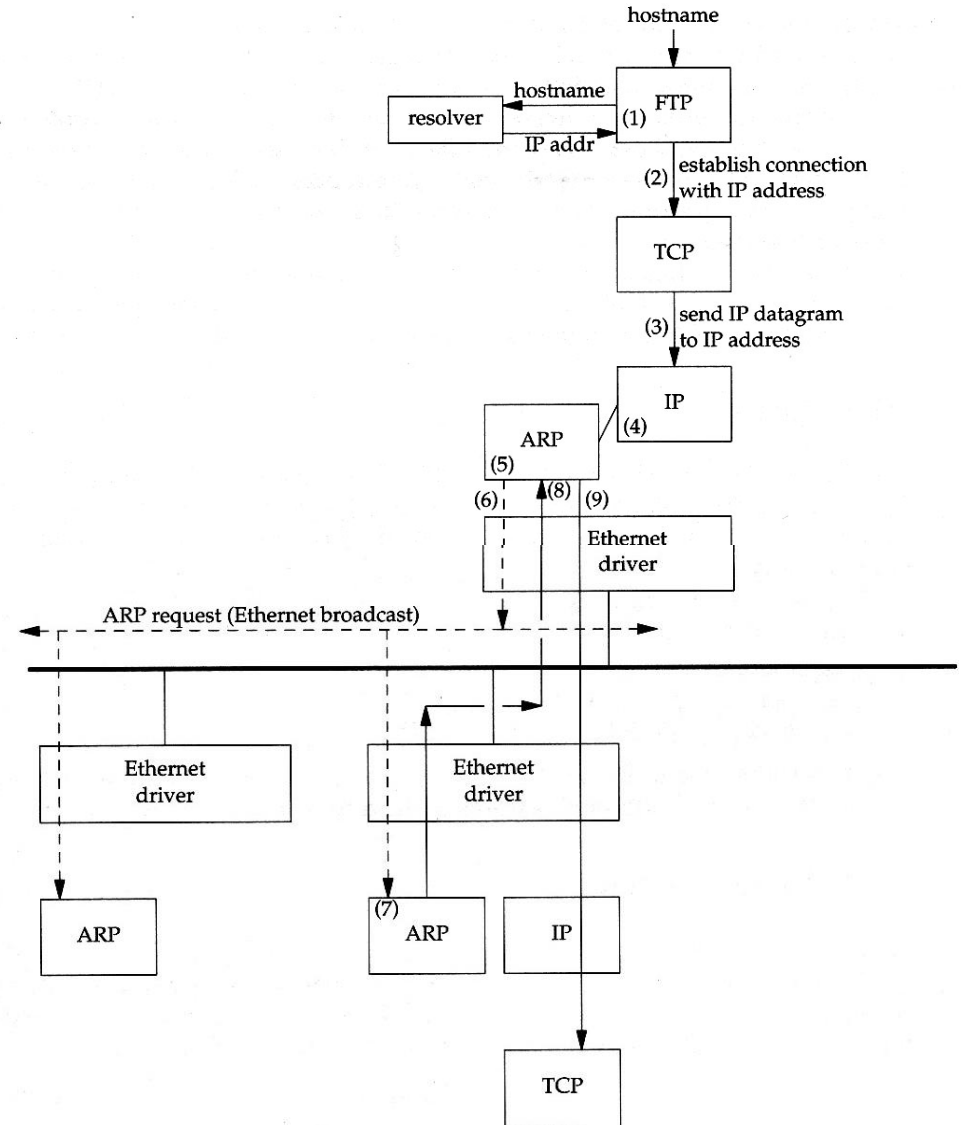
- ◆ A network interface has
  - » one MAC address
  - » one (or more) IP address(es)
- ◆ ARP: Address Resolution Protocol
  - » Protocol used to  
    obtain the MAC address  
    associated to a given IP address
- ◆ RARP – Reverse ARP
  - » Protocol used to  
    obtain the IP address  
    associated to a MAC address




# ARP Example



- **hardware type** : Ethernet=1 ARCNET=7, localtalk=11
- **protocol type** : IP=0x800
- **hlen** : length of hardware address, Ethernet=6 bytes
- **plen** : length of protocol address, IP=4 bytes
- **ARP operation** : ARP request = 1, ARP reply = 2  
RARP request = 3, RARP reply = 4





- ◆ Assume host A sends an IP packet to host B and that this packet is forwarded by router R. What are the MAC and IP addresses (source and destination) observed?
- ◆ What roles does ARP play in this scenario?

---

## *Obtaining IP Addresses*

# *How to Obtain IP Addresses*

---

How does network get subnet part of IP addresss?

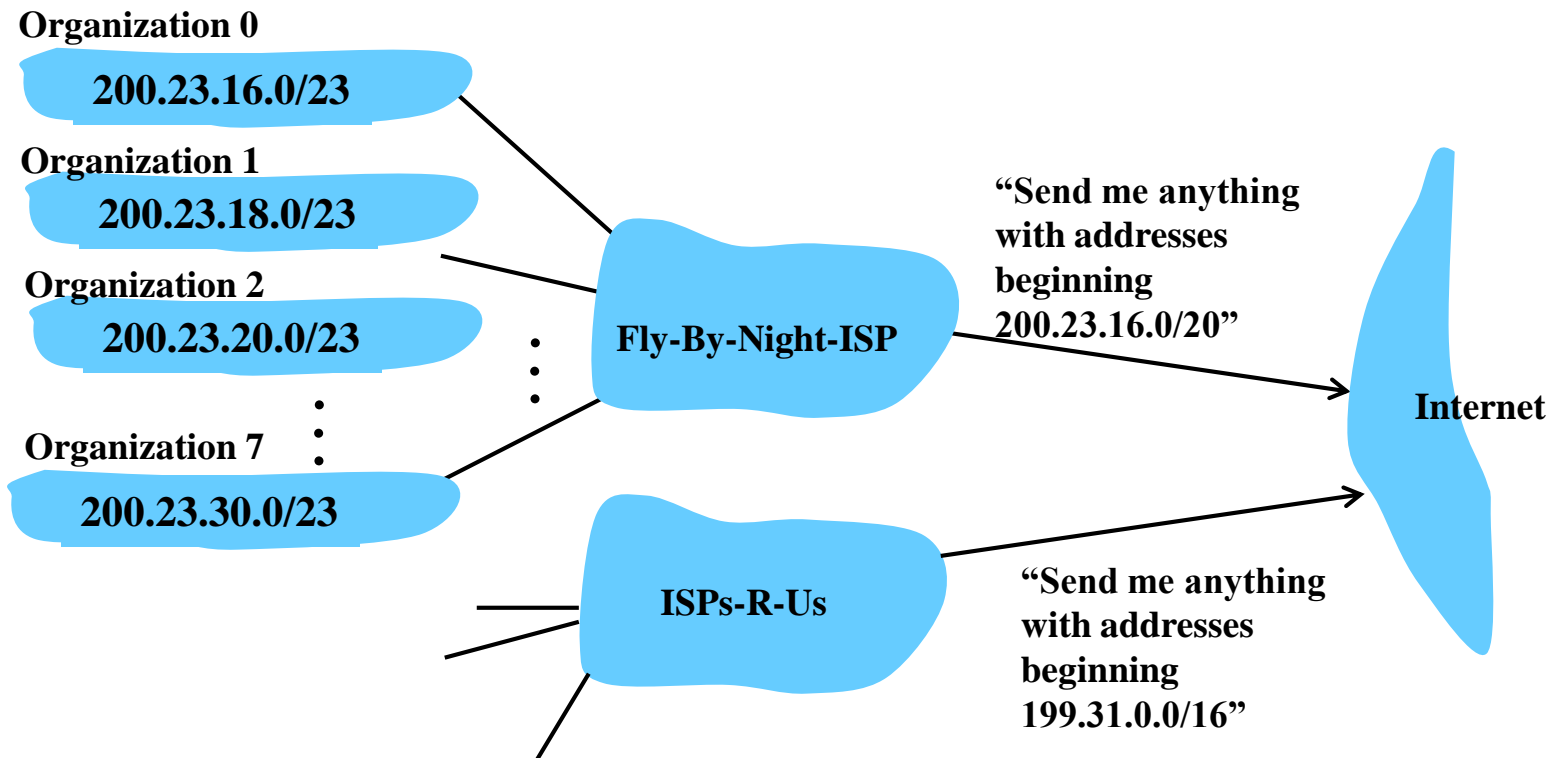
- » Gets allocated portion of its provider ISP's address space

ISP's block	<u>11001000</u>	<u>00010111</u>	<u>00010000</u>	00000000	200.23.16.0/20
Organization 0	<u>11001000</u>	<u>00010111</u>	0001 <u>000</u> 0	00000000	200.23.16.0/23
Organization 1	<u>11001000</u>	<u>00010111</u>	0001 <u>001</u> 0	00000000	200.23.18.0/23
Organization 2	<u>11001000</u>	<u>00010111</u>	0001 <u>010</u> 0	00000000	200.23.20.0/23
...	.....				....
Organization 7	<u>11001000</u>	<u>00010111</u>	0001 <u>111</u> 0	00000000	200.23.30.0/23

# *Hierarchical Addressing - Route Aggregation*

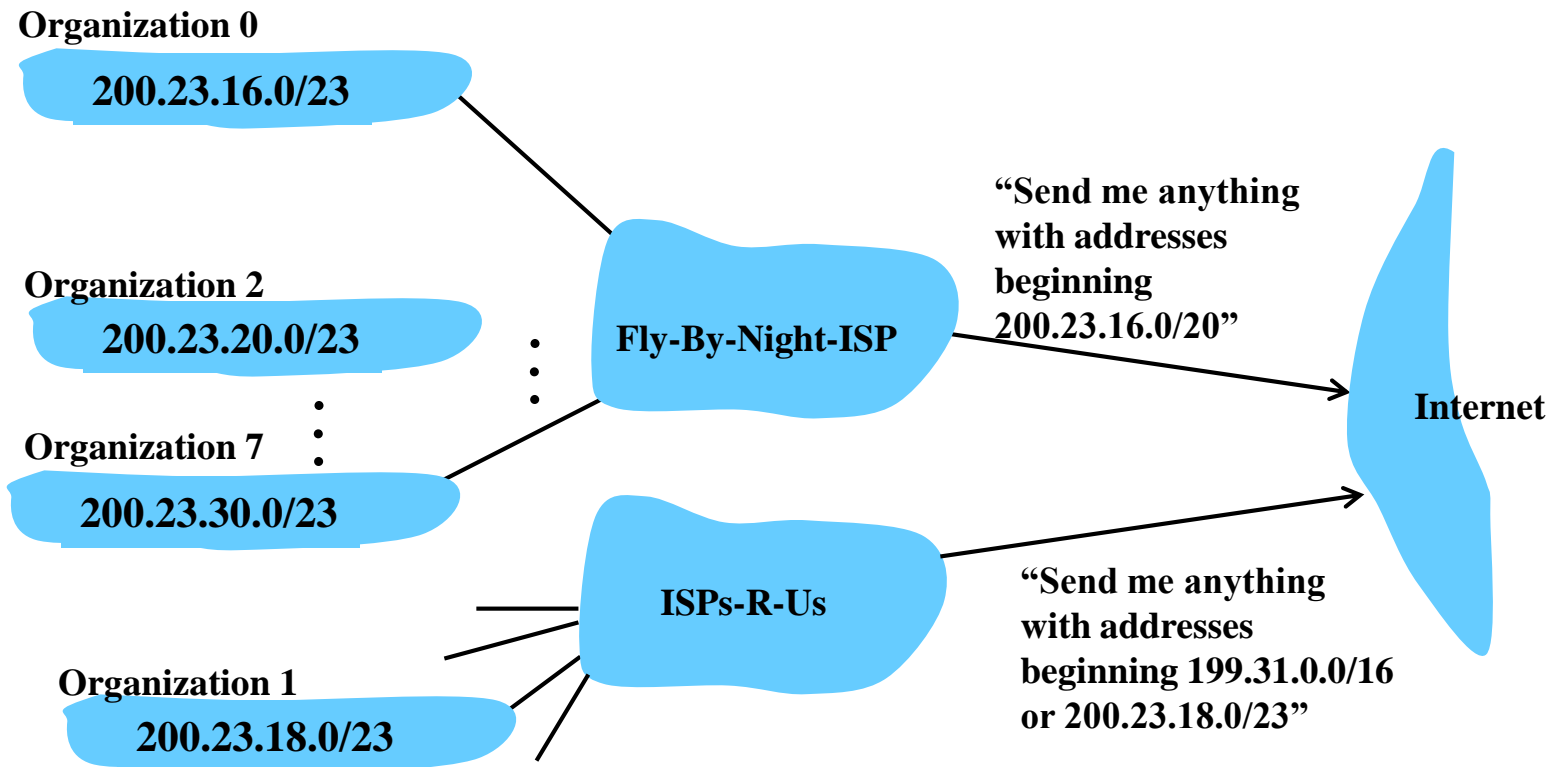
---

- ◆ Hierarchical addressing
  - allows efficient advertisement of routing information



# *Hierarchical Addressing - More Specific Routes*

ISPs-R-Us has a more specific route to Organization 1



# *IP Addressing*

---

How does an ISP get block of addresses?

- » From ICANN: Internet Corporation for Assigned Names and Numbers
- » ICANN
  - allocates addresses
  - manages Domain Name Service (DNS)
  - assigns domain names, resolves disputes

# *IP Addresses*

---

How does a host obtain an IP address?

- » Hard-coded by system admin in a file
  - Windows: control-panel → network → configuration → tcp/ip → properties
  - UNIX: /etc/rc.config, ifconfig
- » DHCP: Dynamic Host Configuration Protocol
  - dynamically get address from a server
  - “plug-and-play”




# *DHCP - Dynamic Host Configuration Protocol*

- ◆ DHCP allows
  - » host to dynamically obtain its IP address from network server when it joins network
  - » reuse of addresses
  
- ◆ DHCP overview
  - » host broadcasts “**DHCP discover**” msg
  - » DHCP server responds with “**DHCP offer**” msg
  - » host requests IP address “**DHCP request**” msg
  - » DHCP server sends address “**DHCP ack**” msg

# *DHCP - Client-server Scenario*

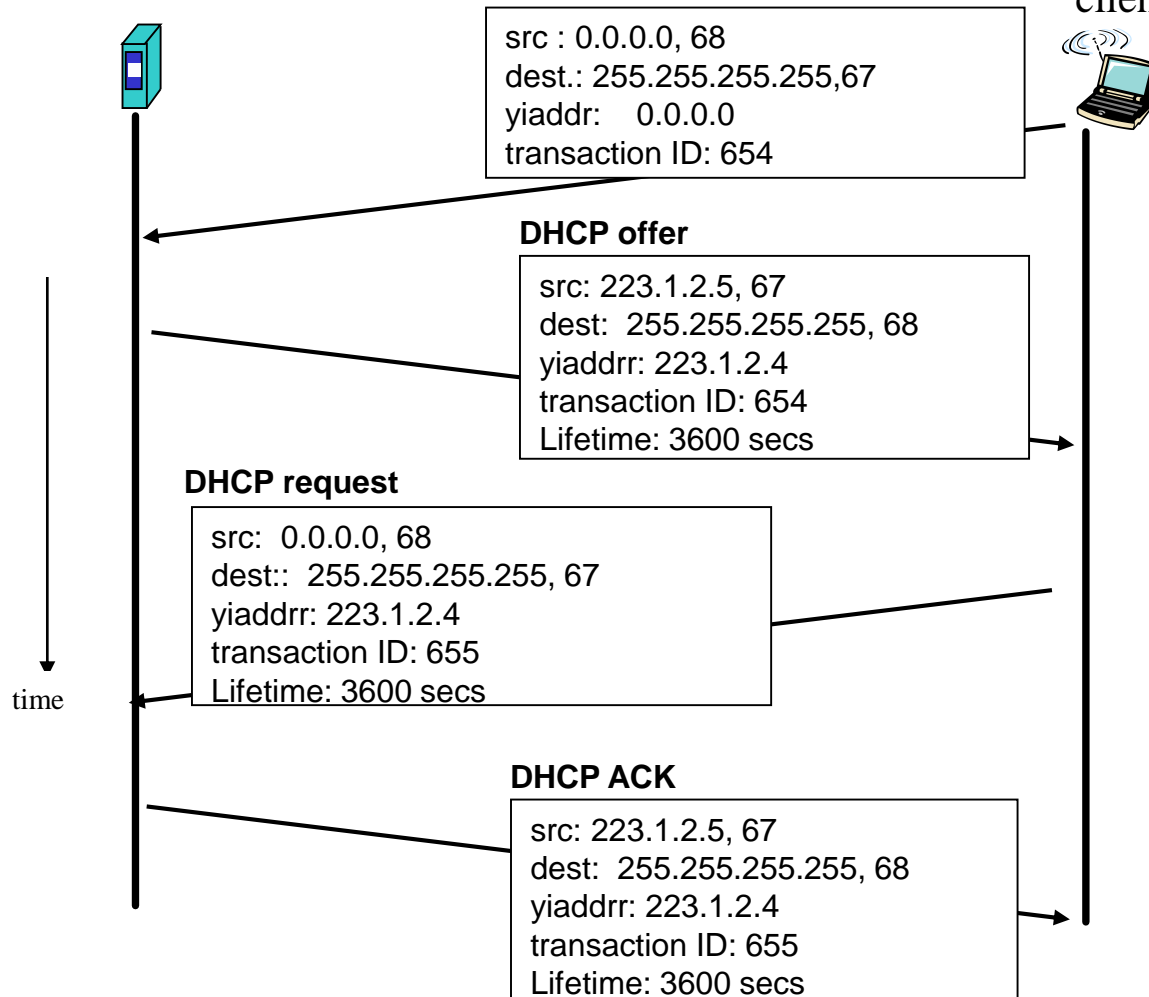
---



# DHCP Client-server

DHCP server: 223.1.2.5

arriving  
client

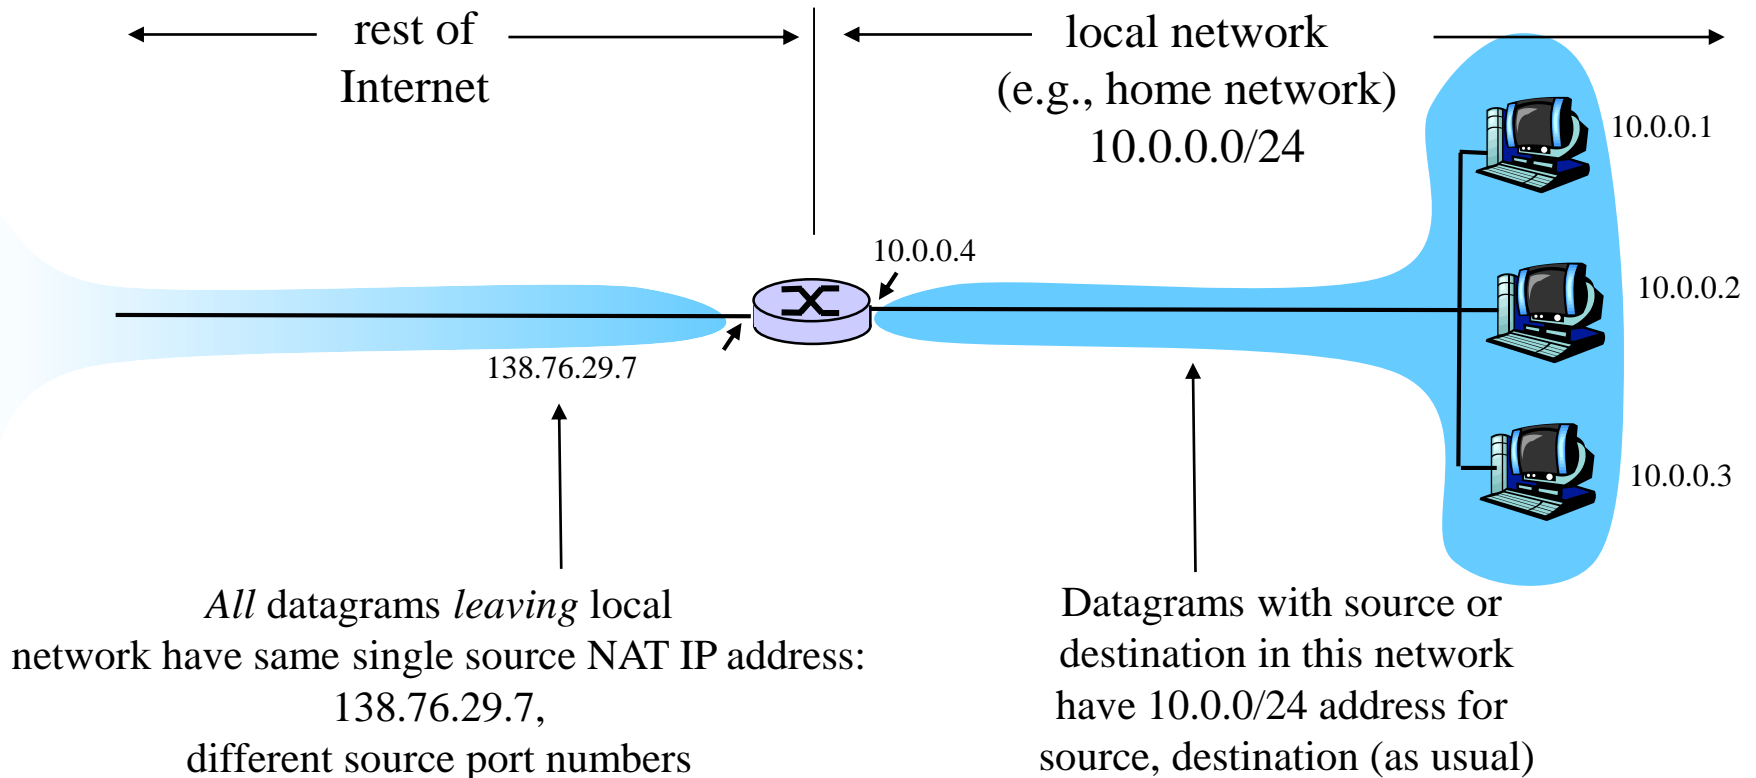


- 
- ♦ Is it sufficient for an arriving client to acquire an IP address?  
What other relevant information shall this client obtain in order to start working with full functionality?

---

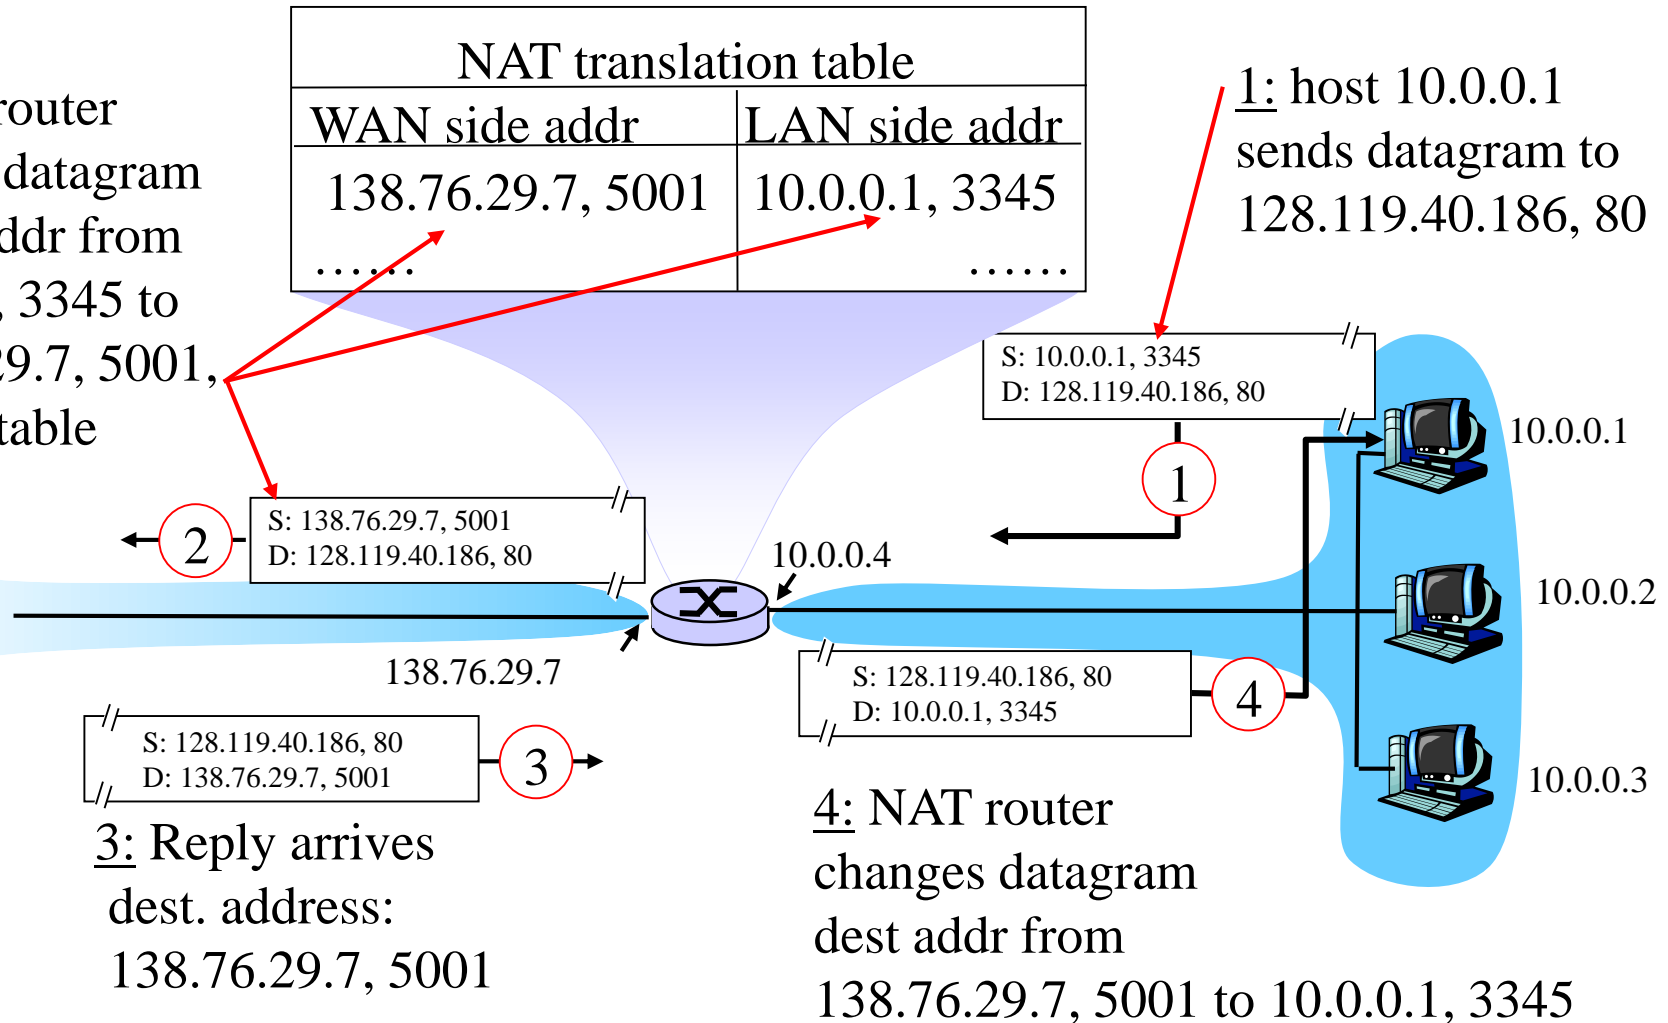
# *Network Address Translation*

# *NAT - Network Address Translation*




# NAT - Network Address Translation

2: NAT router changes datagram source addr from 10.0.0.1, 3345 to 138.76.29.7, 5001, updates table



# NAT Traversal



- ◆ Client wants to connect to server with address 10.0.0.1
  - » but server address 10.0.0.1 is private
  - » only one externally visible NATed address: 138.76.29.7
- ◆ Possible solution – **Port forwarding**
  - » **statically configure NAT**
    - to forward incoming connection requests at given port to server
  - » e.g., (138.76.29.7, port 2500) always forwarded to 10.0.0.1 port 25000



- 
- ◆ How to enable two hosts behind NAT boxes to communicate?  
How does Skype enable voice calls between these two clients?

---

# *Internet Message Control Protocol*

# ICMP - Internet Control Message Protocol

- ♦ Used by router or host
  - » to send layer 3 error or control messages
  - » to other hosts or routers
- ♦ Carried in IP datagrams

0	8	16	31
Type	Code	Checksum	
Unused			
IP Header + 64 bits of original datagram			

(a) Destination Unreachable; Time Exceeded; Source Quench

0	8	16	31
Type	Code	Checksum	
Pointer	Unused		
IP Header + 64 bits of original datagram			

(b) Parameter Problem

0	8	16	31
Type	Code	Checksum	
Gateway Internet Address			
IP Header + 64 bits of original datagram			

(c) Redirect

0		8		16		31	
Type		Code		Checksum			
Identifier				Sequence Number			
Optional data							

(d) Echo, Echo Reply

0	8	16	31
Type	Code	Checksum	
Identifier		Sequence Number	
Originate Timestamp			

(e) Timestamp

0	8	16	31
Type	Code	Checksum	
Identifier		Sequence Number	
Originate Timestamp			
Receive Timestamp			
Transmit Timestamp			

(f) Timestamp Reply

0	8	16	31
Type	Code	Checksum	
Identifier		Sequence Number	

(g) Address Mask Request

0	8	16	31
Type	Code	Checksum	
Identifier		Sequence Number	
Address Mask			

(h) Address Mask Reply

Type	Code	Description
0	0	echo reply (ping)
3	0	dest. network unreachable
3	1	dest host unreachable
3	2	dest protocol unreachable
3	3	dest port unreachable
3	6	dest network unknown
3	7	dest host unknown
4	0	source quench (congestion control - not used)
5		Redirect
8	0	echo request (ping)
9	0	route advertisement
10	0	router discovery
11	0	TTL expired
12	0	bad IP header

# *Ping – Echo Request, Echo Reply*

---

```
sun% ping gemini
PING gemini: 56 data bytes
64 bytes from gemini (140.252.1.11): icmp_seq=0. time=373. ms
64 bytes from gemini (140.252.1.11): icmp_seq=1. time=360. ms
64 bytes from gemini (140.252.1.11): icmp_seq=2. time=340. ms
64 bytes from gemini (140.252.1.11): icmp_seq=3. time=320. ms
64 bytes from gemini (140.252.1.11): icmp_seq=4. time=330. ms
64 bytes from gemini (140.252.1.11): icmp_seq=5. time=310. ms
64 bytes from gemini (140.252.1.11): icmp_seq=6. time=290. ms
64 bytes from gemini (140.252.1.11): icmp_seq=7. time=300. ms
64 bytes from gemini (140.252.1.11): icmp_seq=8. time=280. ms
64 bytes from gemini (140.252.1.11): icmp_seq=9. time=290. ms
64 bytes from gemini (140.252.1.11): icmp_seq=10. time=300. ms
64 bytes from gemini (140.252.1.11): icmp_seq=11. time=280. ms
--gemini PING Statistics--
12 packets transmitted, 12 packets received, 0% packet loss
round-trip (ms) min/avg/max = 280/314/373
```

# Traceroute and ICMP

- ◆ Source sends series of UDP segments to destination

- » first segment has TTL = 1
- » second segment has TTL=2, ...
- » **unlikely port number**


```
svr4% traceroute slip
traceroute to slip (140.252.13.65), 30 hops max. 40 byte packets
1 bsdi (140.252.13.35) 20 ms 10 ms 10 ms
2 slip (140.252.13.65) 120 ms 120 ms 120 ms
```

- ◆ When  $n^{\text{th}}$  datagram arrives

to  $n^{\text{th}}$  router

- » router discards datagram
- » sends to source:  
ICMP TTL expired
- » message includes  
router name & IP address

```
slip% traceroute svr4
traceroute to svr4 (140.252.13.34), 30 hops max, 40 byte packets
1 bsdi (140.252.13.66) 110 ms 110 ms 110 ms
2 svr4 (140.252.13.34) 110 ms 120 ms 110 ms
```



- ◆ When ICMP message arrives, source calculates RTT

- ◆ Traceroute does this 3 times for each TTL

- ◆ Stop criterion

- » UDP segment eventually arrives at destination host
- » Destination returns ICMP “dest port unreachable” packet
- » source stops

# ICMP Redirect

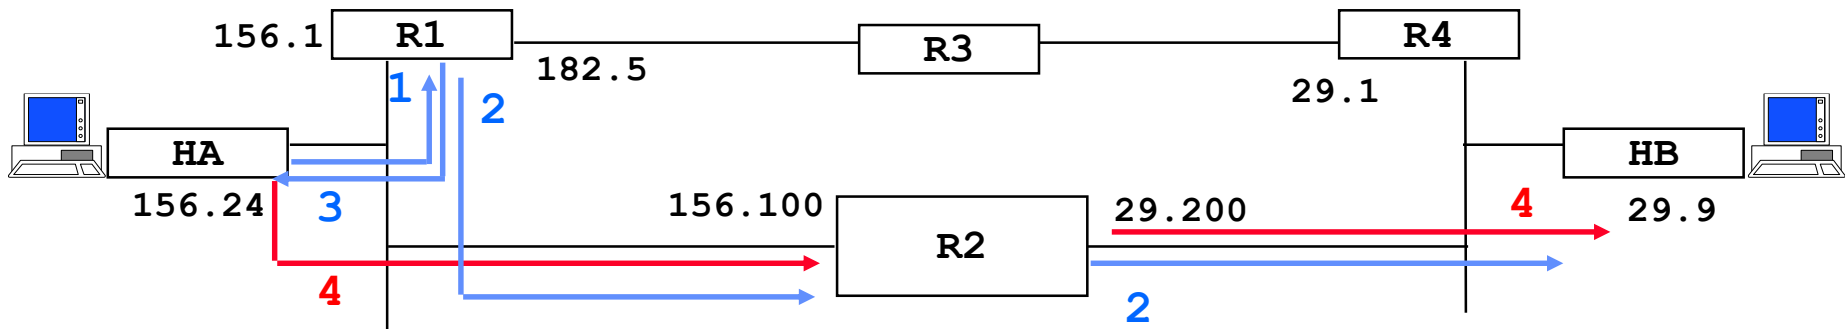
---

- ♦ General routing principle of the TCP/IP architecture
  - » routers have extensive knowledge of routes
  - » hosts have minimal routing information → learn routes also from ICMP redirects
- ♦ ICMP redirect message
  - » Sent by router R1 to source host A
    - when R1 receives a packet from A with destination = B, and R1
      - finds that the next hop is R2 and
      - A is on-link with R2
  - » R1 sends ICMP redirect to A saying next hop for destination B is R2
  - » A updates its forwarding table with a host route

## ICMP Redirect Format

```
/
|
|          IP datagram header  (prot = ICMP)          |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|  Type=5      |      code      |      checksum      |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|
|          Router IP address that should be preferred          |
+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+---+
|
|          IP header plus 8 bytes of original datagram data      |
/
/
```

# ICMP Redirect Example



	dest IP addr	srce IP addr	prot	data part
1:	193.154.29.9	193.154.156.24	udp	xxxxxxx
2:	193.154.29.9	193.154.156.24	udp	xxxxxxx
3:	193.154.156.24	193.154.156.1	icmp	type=redir code=host cksum 193.154.156.100 xxxxxxx (28 bytes of 1)
4:	193.154.29.9	193.154.156.24	udp	.....

# ICMP Redirect Example

---

After 4

```
HA$ netstat -nr
```

Routing Table:

Destination	Gateway	Flags	Interface
-----	-----	-----	-----
127.0.0.1	127.0.0.1	UH	lo0
193.154.29.9	193.154.156.100	UGH	eth0
193.154.156.0	193.154.156.24	U	eth0
224.0.0.0	193.154.156.24	U	eth0
default	193.154.156.1	UG	eth0

Flags:

U - route Up

G - route to a Gateway (next hop router)

H - route to a Host



---

*IPv6*

# *The Need of a New IP*

---

## ◆ IPv4

- Small addressing space (32 bits)
- Non-continuous usage
- Some solutions used to overcome these problems  
private networks (NAT), classless networks (CDIR)

## ◆ IETF developed new IP version: **IPv6**

- Same principles of IPv4
- Many improvements
- Header re-defined

# *IPv6 – Improvements*

---

- » 128 bit addresses (16 octets, 8 *shorts* ). No classes
- » Better QoS support (native flow label)
- » Native security functions (peer authentication, data encryption)
- » Autoconfiguration (*Plug-n-play*)
- » Routing
- » Multicast

# *Address Representation*

---

- ◆ 8 x 16 bit, hexadecimal. Separated by :

**47CD : 1234 : 3200 : 0000 : 0000 : 4325 : B792 : 0428**

- ◆ Compressed format: **FF01:0:0:0:0:0:0:43 → FF01::43**

- ◆ Compatibility with IPv4: **0:0:0:0:0:0:13.1.68.3 or ::13.1.68.3**

- ◆ Loopback address: **::1**

- ◆ Network prefix described by / , same as IPv4
  - » **FEDC:BA98:7600::/40 → network prefix = 40 bits**

# *Reserved Addresses*

---

Allocation	Prefix (binary)	Fraction of Address Space
-----	-----	-----
Unassigned	0000 0000	1/256
Unassigned	0000 0001	1/256
Reserved for NSAP Allocation	0000 001	1/128
Unassigned	0000 01	1/64
Unassigned	0000 1	1/32
Unassigned	0001	1/16
Global Unicast	001	1/8
Unassigned	010	1/8
Unassigned	011	1/8
Unassigned	100	1/8
Unassigned	101	1/8
Unassigned	110	1/8
Unassigned	1110	1/16
Unassigned	1111 0	1/32
Unassigned	1111 10	1/64
Unassigned	1111 110	1/128
Unassigned	1111 1110 0	1/512
Link-Local Unicast Addresses	1111 1110 10	1/1024
Site-Local Unicast Addresses	1111 1110 11	1/1024
Multicast Addresses	1111 1111	1/256

# *Adresses – Link-Local, Global Unicast, Anycast, Multicast*

---

## » *Link-Local*

- Used for communication between hosts in the same LAN/link
- Address built from MAC address
- Routers do not forward packets having *Link-Local* destination addresses

## » *Global Unicast*

- Global addresses
- Address: network prefix + computer identifier
- Structured prefixes

Network aggregation; less entries in the router forwarding tables

## » *Anycast*

- Group address; packet is received by any (only one) member of the group

## » *Multicast*

- Group address; packet received by all the members of the group

# Address Formats

---

n bits	m bits	128-n-m bits
001 global rout prefix	subnet ID	interface ID

Global Unicast Address  
(2000::/3)

10 bits	54 bits	64 bits
1111111010	0	interface ID

Link-Local Unicast address  
(fe80::/10)

10 bits	54 bits	64 bits
1111111011	subnet ID	interface ID

Site-Local Unicast address  
(fec0::/10) (not used)

n bits	128-n bits
subnet prefix	0000000000000000

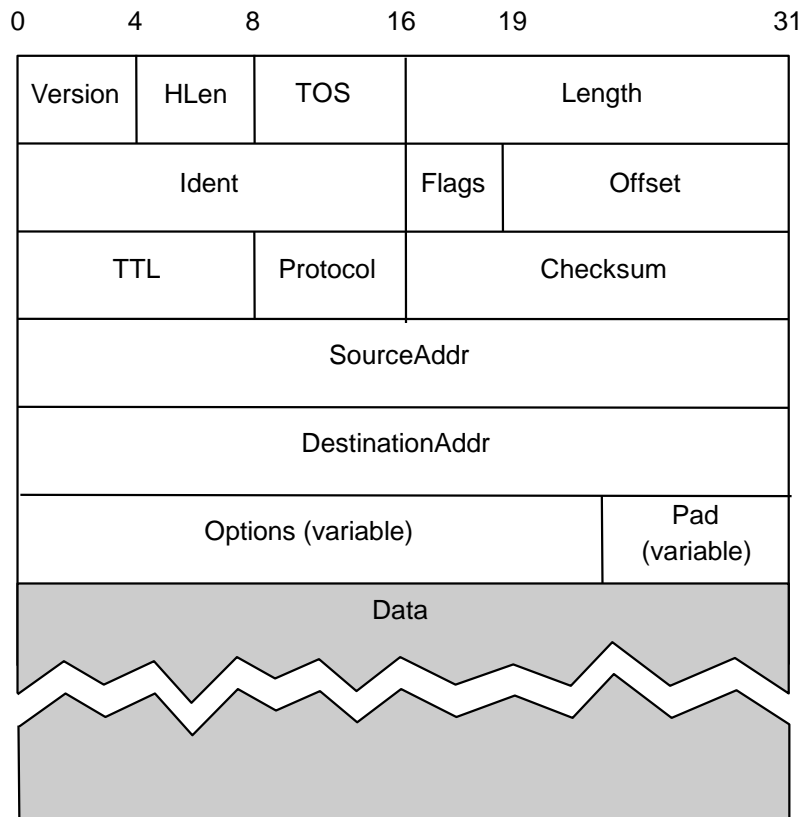
Anycast address

8	4	4	112 bits
11111111	flgs	scop	group ID

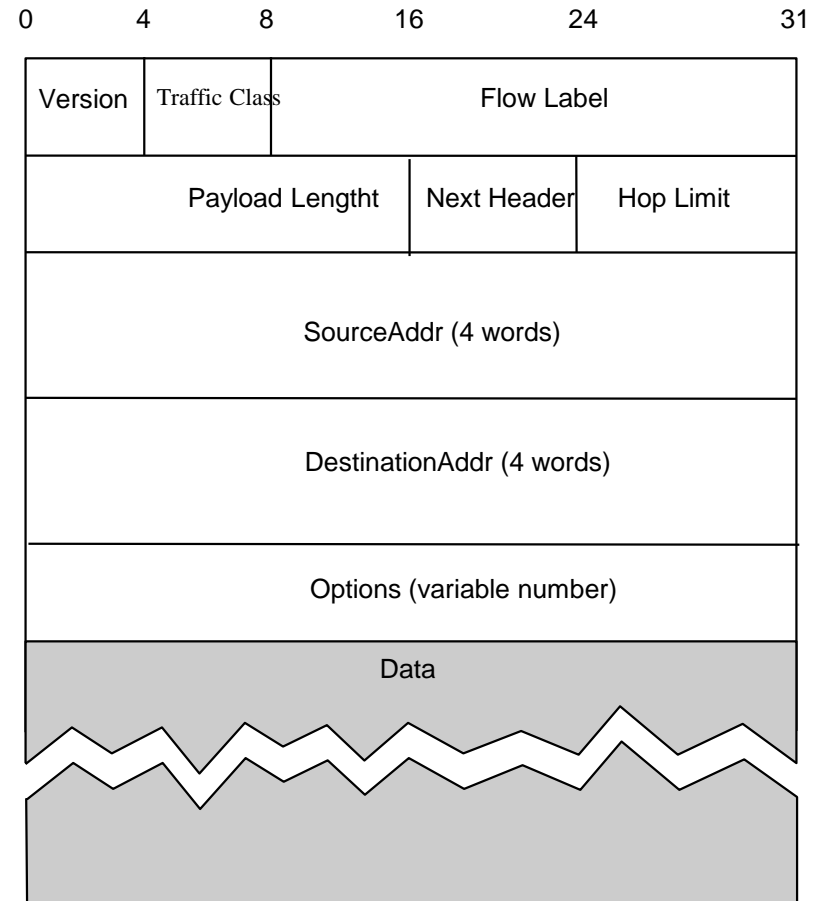
Multicast address  
Scope - link, site, global, ...  
(ff::/8)

# Headers IPv4 and IPv6

---



IPv4

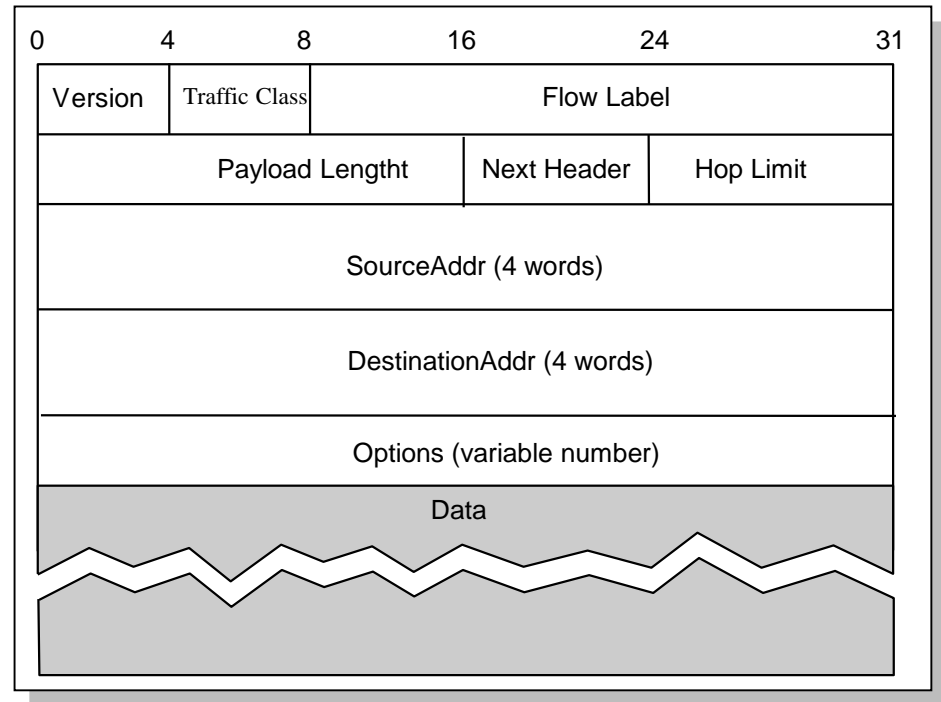


IPv6

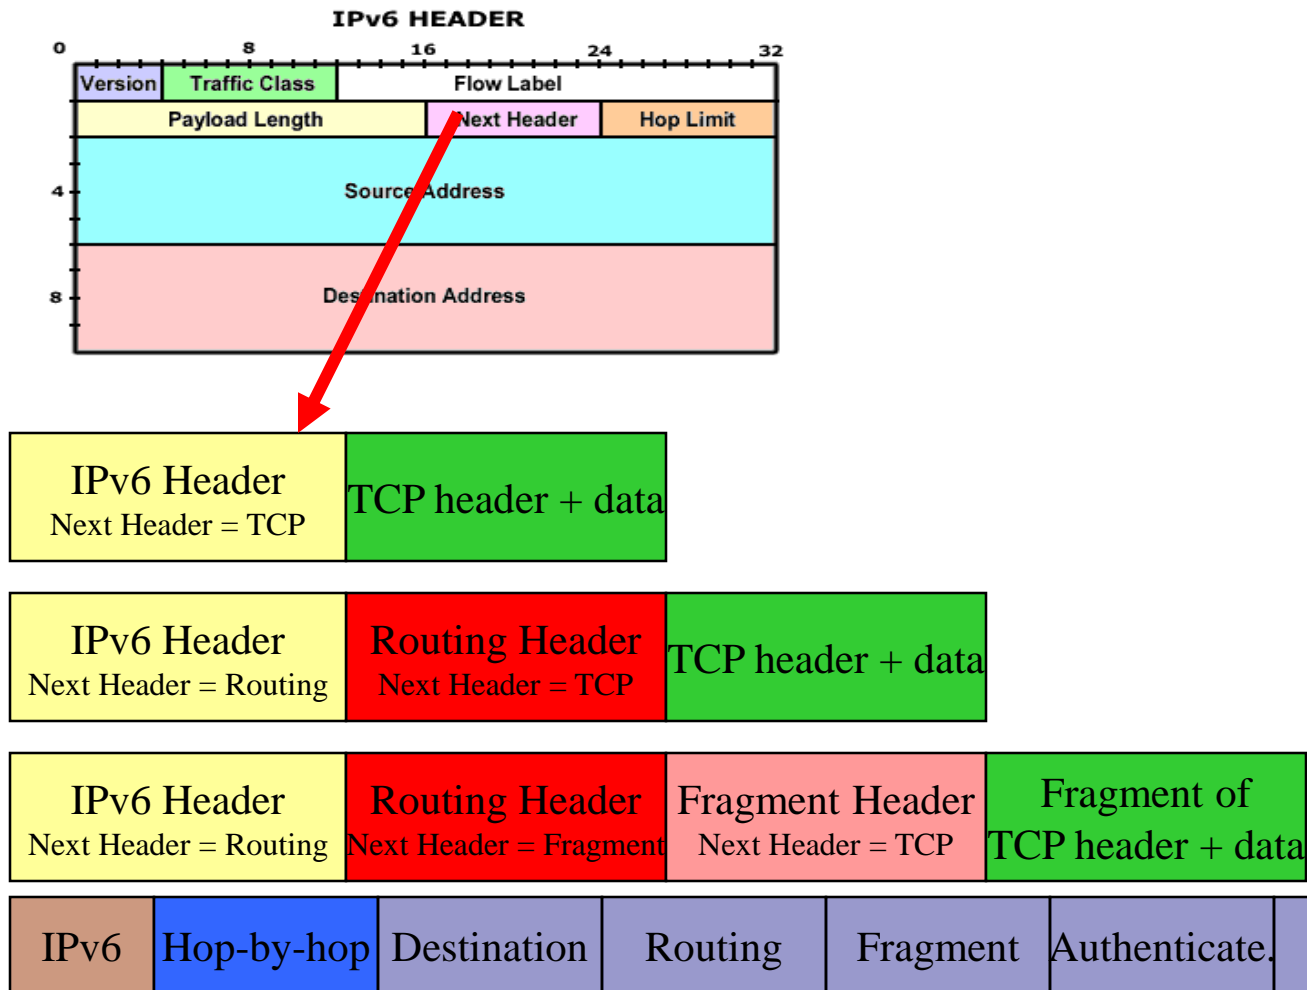


# IPv6 Header

- ♦ Flow label → identifies packet flow
  - » QoS, resource reservation
  - » Packets receive same service
- ♦ Payload length
  - » Header not included
- ♦ Hop limit = TTL (v4)
- ♦ Next header
  - » Identifies next header/extension
- ♦ Options → included as extension headers



# Extension Headers



# *Extension Headers*

---

- » Hop-by-hop

additional information, inspected by every node traversed by the packet

Other headers are inspected only at the destination or at pre-defined nodes

- » Destination: Information for the destination node


- » Routing: List of nodes to be visited by the packet

- » Fragmentation: Made by the source; it shall find MPU

- » Authentication: Authentication (signature) of packet header

- » ESP: Data encryption

# Example of Lab Network



# Configuration examples in Linux

```
tux13:~# /sbin/ifconfig eth0 inet6 add 2000:0:0:1::1/64
tux13:~# ifconfig eth0
eth0      Link encap:Ethernet  HWaddr 00:C0:DF:08:D5:99
          inet addr:172.16.1.13  Bcast:172.16.1.255  Mask:255.255.255.0
          inet6 addr: 2000:0:0:1::1/64 Scope:Global
          inet6 addr: fe80::2c0:dfff:fe08:d599/10 Scope:Link
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:81403 errors:0 dropped:0 overruns:0 frame:0
          TX packets:2429 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:100
          RX bytes:4981344 (4.7 MiB)  TX bytes:260692 (254.5 KiB)
          Interrupt:5
```

```
tux13:~# /sbin/route -A inet6 add 2000::/3 gw 2000:0:0:1::aa
```

```
tux13:~# route -A inet6
```

Kernel IPv6 routing table

Destination	NextHop	Flags	Metric	Ref	Use	Iface
::1/128	::	U	0	0	0	lo
2000:0:0:1::1/128	::	U	0	0	0	lo
2000:0:0:1::/64	::	UA	256	0	0	eth0
2000::/3	2000:0:0:1::aa	UG	1	0	0	eth0
fe80::2c0:dfff:fe08:d599/128	::	U	0	0	0	lo
fe80::/10	::	UA	256	0	0	eth0
ff00::/8	::	UA	256	0	0	eth0
::/0	::	UDA	256	0	0	eth0

# *Identifier IEEE EUI-64*

---

Method to create a IEEE EUI-64 identifier from an IEEE 48bit MAC identifier.  
This is to insert two octets, with hexadecimal values of 0xFF and 0xFE,  
in the middle of the 48 bit MAC (between the company\_id and vendor supplied id).  
For example, the 48 bit IEEE MAC with global scope:

0	1   1	3   3	4
0	5   6	1   2	7
+-----+-----+-----+			
ccccc0gcccccccc   cccccccmmmmmmmmmm   mmmmmmmmmmmmmmmmmmm			
+-----+-----+-----+ 00:C0:DF:08:D5:99			

where "c" are the bits of the assigned company\_id, "0" is the value of the universal/local bit to indicate global scope, "g" is individual/group bit, and "m" are the bits of the manufacturer-selected extension identifier.

The interface identifier would be of the form:

0	1   1	3   3	4   4	6
0	5   6	1   2	7   8	3
+-----+-----+-----+-----+				
cccccclgcccccccc   cccccccc1111111   11111110mmmmmmmmmm   mmmmmmmmmmmmmmmmmmm				
+-----+-----+-----+-----+				

fe80::2c0:dfff:fe08:d599

# *Protocol Neighbor Discovery (ND)*

---

- ◆ IPv6 node uses ND for
  - » Find other nodes in the same link /LAN
  - » Find a node MAC address
    - ND substitutes ARP
  - » Find router(s) in its network
  - » Maintaining information about neighbour nodes
  
- ◆ ND similar to the IPv4 functions
  - » ARP IPv4
  - » ICMP Router Discovery
  - » ICMP Redirect

# *ND Messages*

---

- » ICMP messages (over IP); using *Link Local* addresses
- » ***Neighbor Solicitation***  
Sent by a host to obtain MAC address of a neighbour / to verify its presence
- » ***Neighbor Advertisement:*** Answer to the request
- » ***Router Advertisement***  
Information about the network prefix; periodic or under request  
Sent by router to IP address *Link Local multicast*
- » ***Router Solicitation:*** host solicits from router a *Router Advertisement* message
- » ***Redirect:*** Used by a router to inform a host about the best route to a destination



# *Homework*

---

1. Review slides
2. Read from Kurose&Ross
  - » Chapter 4 – The Network Layer  
(this set of slides follows mainly Kurose&Ross)
3. Or, from Tanenbaum,
  - » Chapter 5 – The Network Layer
4. Answer questions at moodle