

---

*Redes de Computadores*

# **Shortest Paths in Networks**


*Manuel P. Ricardo*

*Faculdade de Engenharia da Universidade do Porto*


- 
- » *What is a graph?*
  - » *What is a spanning tree?*
  - » *What is a shortest path tree?*
  - » *How are paths defined in a network?*
  - » *How does the Dijkstra algorithm work?*
  - » *How does a link state routing protocol work?*
  - » *How does a node learn about neighbours?*
  - » *How does the Bellman-Ford algorithm work?*
  - » *How does a distance vector work?*
  - » *What are the limitations of the layer 2 network of switches?*
  - » *How does the IEEE spanning tree protocol work?*
  - » *What is the maximum capacity of a flow network?*

# *Graph – Directed and Undirected*

---



**a) Directed graph**



**b) Undirected graph**

$$G = (V, E)$$

$$V = \{v_1, v_2, v_3, v_4, v_5, v_6\}, \quad |V| = 6$$

$$E = \{(v_1, v_2), (v_2, v_3), (v_2, v_4), (v_3, v_4), \\ (v_4, v_3), (v_3, v_5), (v_4, v_5), (v_5, v_6)\}, \quad |E| = 8$$

$$G = (V, E)$$


$$V = \{v_1, v_2, v_3, v_4, v_5, v_6\}, \quad |V| = 6$$

$$E = \{(v_1, v_2), (v_2, v_3), (v_2, v_4), (v_3, v_4), \\ (v_3, v_5), (v_4, v_5), (v_5, v_6)\}, \quad |E| = 7$$

# Tree

---

- ◆ Trees  $T = (V, E)$ 
  - » graph with no cycles
  - » number of edges  $|E| = |V| - 1$
  - » any two vertices of the tree are connected by exactly one path
- ◆ A tree  $T$  is said to span a graph  $G = (V, E)$  (spanning tree) if
  - »  $T = (V, E')$  and  $E' \subseteq E$




# Shortest Path Trees

---

- ♦ Graphs and Trees can be weighted

- »  $G=(V,E,w)$
- »  $T=(V, E',w)$
- »  $w: E \rightarrow \mathbb{R}$




- ♦ Total cost of a tree  $T \rightarrow C_{total}(T) = \sum_{i=1}^{|E|} w(e_i)$
- ♦ Minimum Spanning Tree  $T^* \rightarrow C_{total}(T^*) = \min(C_{total}(T))$ 
  - » algorithms used to compute MST: Prim, Kruskal
- ♦ **Shortest Path Tree (SPT) Rooted at Vertex s**
  - » tree composed by the
  - » **union of the shortest paths between s and each of other vertices of G**
  - » Algorithms used to compute SPT: **Dijkstra, Bellman-Ford**
- ♦ Computer networks use **Shortest Path Trees**

---

## *Routing in Layer 3 Networks*

# Forwarding, Routing

- ♦ Forwarding → data plane
  - » directing packet from input to output link
  - » using a forwarding table
- ♦ Routing → control plane
  - » computing paths the packets will follow
  - » routers exchange messages
  - » each router creates its forwarding table



# *Importance of Routing*

---

- ◆ End-to-end performance
  - » path affects quality of service
  - » delay, throughput, packet loss
  
- ◆ Use of network resources
  - » balance traffic over routers and links
  - » avoiding congestion by directing traffic to less-loaded links
  
- ◆ Transient disruptions
  - » failures, maintenance
  - » limiting packet loss and delay during changes




# *Shortest-Path Routing*

---

## Path-selection model


- » Destination-based
- » Load-insensitive (e.g., static link weights)
- » Minimum hop count or minimum sum of link weights



# Shortest-Path Problem

---

- ◆ Given a network topology with link costs
  - »  $c(x,y)$  - link cost from node  $x$  to node  $y$
  - » infinity if  $x$  and  $y$  are not direct neighbors
- ◆ Compute the least-cost paths from source  $u$  to all nodes
  - $p(v)$  - node predecessor of node  $v$  in the path to  $u$



# *Dijkstra's Shortest-Path Algorithm*

---

- ♦ Iterative algorithm
  - » After  $k$  iterations  $\rightarrow$  known least-cost paths to  $k$  nodes
- ♦  $\mathbf{S} \rightarrow$  set of nodes for which least-cost path is known
  - » Initially,  $\mathbf{S} = \{\mathbf{u}\}$ , where  $\mathbf{u}$  is the source node
  - » Add one node to  $\mathbf{S}$  in each iteration
- ♦  $\mathbf{D}(\mathbf{v}) \rightarrow$  current cost of path from source to node  $\mathbf{v}$ 
  - » Initially
    - $\mathbf{D}(\mathbf{v}) = \mathbf{c}(\mathbf{u}, \mathbf{v})$  for all nodes  $\mathbf{v}$  adjacent to  $\mathbf{u}$
    - $\mathbf{D}(\mathbf{v}) = \infty$  for all other nodes  $\mathbf{v}$
  - » Continually update  $\mathbf{D}(\mathbf{v})$  when shorter paths are learned

# *Dijkstra's Algorithm*

---

1 *Initialization:*

2  $S = \{u\}$

3 for all nodes  $v$

4 if  $v$  adjacent to  $u$  {

5  $D(v) = c(u,v)$  }

6 else  $D(v) = \infty$

7

8 *Loop*

9 find node  $w$  not in  $S$  with the smallest  $D(w)$

10 add  $w$  to  $S$





11 update  $D(v)$  for all  $v$  adjacent to  $w$  and not in  $S$ :

12  $D(v) = \min\{D(v), D(w) + c(w,v)\}$

13 *until all nodes in  $S$*





# *Dijkstra's Algorithm - Example*

---




# *Dijkstra's Algorithm - Example*

---



# Shortest-Path Tree

◆ Shortest-path tree from u



◆ Forwarding table at u

	link
v	(u,v)
w	(u,w)
x	(u,w)
y	(u,v)
z	(u,v)
s	(u,w)
t	(u,w)

# *Link-State Routing*

---


- ◆ Each router keeps track of its incident links
  - » link up, link down
  - » cost on the link
- ◆ Each router broadcasts link state
  - every router gets a complete view of the graph
- ◆ Each router runs Dijkstra's algorithm, to
  - » compute the shortest paths
  - » construct the forwarding table
- ◆ Example protocols
  - » Open Shortest Path First (OSPF)
  - » Intermediate System – Intermediate System (IS-IS)



# *Detection of Topology Changes*

---





- ♦ Beacons generated by routers on links
  - » Periodic “hello” messages in both directions
  - » few missed “hellos” → link failure



# *Broadcasting the Link State*

---


- ◆ How to Flood the link state?
  - » every node sends link-state information through adjacent links
  - » next nodes forward that info to all links except the one where the information arrived
  
- ◆ When to initiate flooding?
  - » Topology change
    - link or node failure/recovery
    - link cost change
  - » Periodically
    - refresh link-state information
    - typically 30 minutes



# *Scaling Link-State Routing*

---


- ◆ Overhead of link-state routing
  - » flooding link-state packets throughout the network
  - » running Dijkstra's shortest-path algorithm
- ◆ Introducing hierarchy through “areas”



# *Bellman-Ford Algorithm*

---

- ◆ Define distances at each node  $x$ 
  - »  $d_x(y) = \text{cost of least-cost path from } x \text{ to } y$
- ◆ Update distances based on neighbors
  - »  $d_x(y) = \min \{c(x,v) + d_v(y)\}$  over all neighbors  $v$



$$d_u(z) = \min\{c(u,v) + d_v(z), \\ c(u,w) + d_w(z)\}$$

# *Distance Vector Algorithm*

---

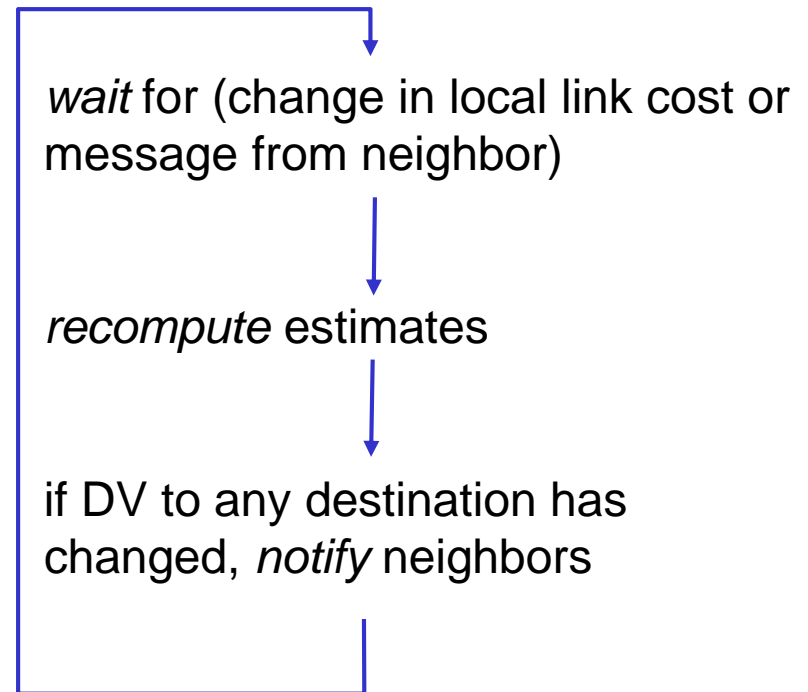
- ♦  $c(x,v)$  = cost for direct link from  $x$  to  $v$   
node  $x$  maintains costs of direct links  $c(x,v)$
- ♦  $D_x(y)$  = estimate of least cost from  $x$  to  $y$   
node  $x$  maintains distance vector  $\mathbf{D}_x = [D_x(y): y \in N]$
- ♦ Node  $x$  maintains also its neighbors' distance vectors  
for each neighbor  $v$ ,  $x$  maintains  $\mathbf{D}_v = [D_v(y): y \in N]$
- ♦ Each node  $v$  periodically sends  $\mathbf{D}_v$  to its neighbors
  - » and neighbors update their own distance vectors
  - »  $D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\}$  for each node  $y \in N$
- ♦ Over time, the distance vector  $\mathbf{D}_x$  converges

# *Distance Vector Algorithm*

---

- ◆ Iterative, asynchronous
  - each local iteration caused by:
    - local link cost change
    - distance vector update message from neighbor
- ◆ Distributed
  - » node notifies neighbors only when its DV changes
- ◆ Neighbors then notify their neighbors, if necessary

Each node:



# Distance Vector Example - Step 0

Table for A			Table for B		
Dst	Cst	Hop	Dst	Cst	Hop
A	0	A	A	4	A
B	4	B	B	0	B
C	$\infty$	—	C	$\infty$	—
D	$\infty$	—	D	3	D
E	2	E	E	$\infty$	—
F	6	F	F	1	F




Table for C			Table for D			Table for E			Table for F		
Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop
A	$\infty$	—	A	$\infty$	—	A	2	A	A	6	A
B	$\infty$	—	B	3	B	B	$\infty$	—	B	1	B
C	0	C	C	1	C	C	$\infty$	—	C	1	C
D	1	D	D	0	D	D	$\infty$	—	D	$\infty$	—
E	$\infty$	—	E	$\infty$	—	E	0	E	E	3	E
F	1	F	F	$\infty$	—	F	3	F	F	0	F

# Distance Vector Example - Step 1

Table for A			Table for B		
Dst	Cst	Hop	Dst	Cst	Hop
A	0	A	A	4	A
B	4	B	B	0	B
C	7	F	C	2	F
D	7	B	D	3	D
E	2	E	E	4	F
F	5	E	F	1	F




Table for C			Table for D			Table for E			Table for F		
Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop
A	7	F	A	7	B	A	2	A	A	5	B
B	2	F	B	3	B	B	4	F	B	1	B
C	0	C	C	1	C	C	4	F	C	1	C
D	1	D	D	0	D	D	$\infty$	—	D	2	C
E	4	F	E	$\infty$	—	E	0	E	E	3	E
F	1	F	F	2	C	F	3	F	F	0	F



# Distance Vector Example - Step 2

Table for A			Table for B		
Dst	Cst	Hop	Dst	Cst	Hop
A	0	A	A	4	A
B	4	B	B	0	B
C	6	E	C	2	F
D	7	B	D	3	D
E	2	E	E	4	F
F	5	E	F	1	F




Table for C			Table for D			Table for E			Table for F		
Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop	Dst	Cst	Hop
A	6	F	A	7	B	A	2	A	A	5	B
B	2	F	B	3	B	B	4	F	B	1	B
C	0	C	C	1	C	C	4	F	C	1	C
D	1	D	D	0	D	D	5	F	D	2	C
E	4	F	E	5	C	E	0	E	E	3	E
F	1	F	F	2	C	F	3	F	F	0	F


# *Routing Information Protocol (RIP)*

---

- ◆ Distance vector protocol
  - » nodes send distance vectors every 30 seconds
  - » or, when an update causes a change in routing
- ◆ RIP is limited to small networks

# *BGP – The Exterior Gateway Routing Protocol*

---



(a)

Information F receives  
from its neighbors about D

From B: "I use BCD"  
From G: "I use GCD"  
From I: "I use IFGCD"  
From E: "I use EFGCD"

(b)

(a) A set of BGP routers.      (b) Information sent to F

---

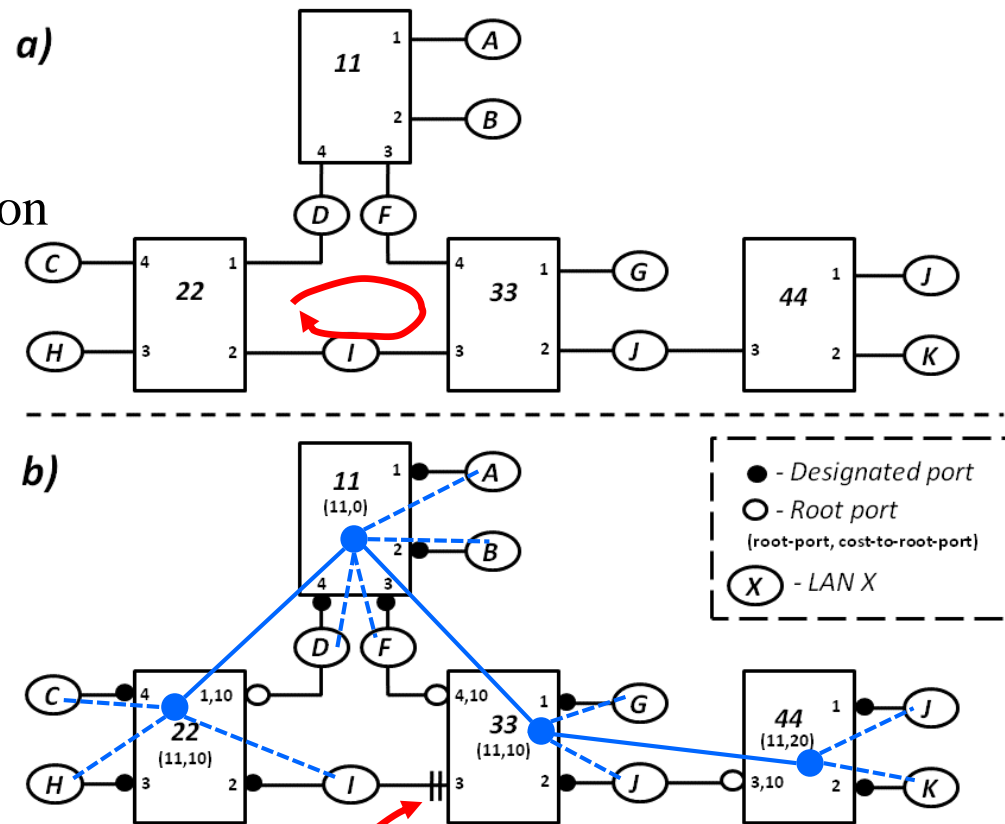
# *Unique Spanning Tree in Ethernet Networks*

# L2 Networking - Single Tree Required

- Ethernet frame
  - No hop-count
  - Could loop forever
  - broadcast frame, mis-configuration

- Layer 2 network
  - **Required to have tree topology**
  - Single path between every pair of stations

- Spanning Tree Protocol (STP)
  - Running in bridges
  - Helps building the spanning tree
  - Blocks ports




# *Constructing a Spanning Tree*

---

## Distributed algorithm

- » switches need to elect a “root”
  - the switch with the smallest identifier
- » each switch identifies if its interface is on **the shortest path from the root**
- » messages (Y, d, X)
  - from node X
  - claiming Y is the root
  - and the distance is d



# *Steps in Spanning Tree Algorithm*


---

- ◆ Initially, each switch thinks it is the root
  - » switch sends a message out every interface
  - » identifying itself as the root with distance 0
  - » example: switch X announces (X, 0, X)
  
- ◆ Other switches update their view of the root
  - » upon receiving a message, check the root id
  - » if the new id is smaller, start viewing that switch as root
  
- ◆ Switches compute their distance from the root
  - » add 1 to the distance received from a neighbor
  - » identify interfaces not on a shortest path to the root and exclude them from the spanning tree

## *Example - Switch #4's Viewpoint*

---

- ◆ Switch #4 thinks it is the root
  - » sends (4, 0, 4) message to 2 and 7
- ◆ Then, switch #4 hears from #2
  - » receives (2, 0, 2) message from 2
  - » ... and thinks that #2 is the root
  - » and realizes it is just one hop away
- ◆ Then, switch #4 hears from #7
  - » receives (2, 1, 7) from 7
  - » and realizes this is a longer path
  - » so, prefers its own one-hop path
  - » and removes 4-7 link from the tree






## *Example - Switch #4's Viewpoint*

---

- ◆ Switch #2 hears about switch #1
  - » switch 2 hears (1, 1, 3) from 3
  - » switch 2 starts treating 1 as root
  - » and sends (1, 2, 2) to neighbors
- ◆ Switch #4 hears from switch #2
  - » switch 4 starts treating 1 as root
  - » and sends (1, 3, 4) to neighbors
- ◆ Switch #4 hears from switch #7
  - » switch 4 receives (1, 3, 7) from 7
  - » and realizes this is a longer path
  - » so, prefers its own three-hop path
  - » and removes 4-7 link from the tree



---

## *Maximum Flow of a Network*


# *Flow Network Model*

---

- ♦ **Flow network**

- » source s
- » sink t
- » nodes a, b and c

- ♦ Edges are labeled with **capacities**
  - » (e.g. bit/s)



- ♦ Communication networks are not flow networks
  - » they are queue networks
  - » flow networks enable to determine limit values

# *Maximum Capacity of a Flow Network*

---





- ♦ Max-flow min-cut theorem
  - » maximum amount of flow transferable through a network
  - » equals minimum value among all simple cuts of the network
- ♦ Cut  $\rightarrow$  split of the nodes  $V$  into two disjoint sets  $S$  and  $T$ 
  - »  $S \cup T = V$
  - » there are  $2^{|V|-2}$  possible cuts
- ♦ Capacity of cut  $(S, T)$ : 
$$c(S, T) = \sum_{(u,v) \mid u \in S, v \in T, (u,v) \in E} c(u, v)$$

# Max-flow Min-cut - Example

$2^{|V|-2} = 8$  possible cuts

Cut	Vertices					$c(S, T)$	Feasibility
	s	a	b	c	t		
1	S	S	S	S	T	10	✓
2	S	S	S	T	T	19	✓
3	S	S	T	S	T	13	✓
4	S	S	T	T	T	17	✓
5	S	T	S	S	T	-	×
6	S	T	S	T	T	-	×
7	S	T	T	S	T	11	✓
8	S	T	T	T	T	12	✓

Maximum flow = 10



# *Homework*

---

1. Review slides
2. Read from Tanenbaum
  - » Section 5.2 – Routing algorithms
  - » Section 4.8.3 Spanning Tree Bridges