

# Chapter 2

## Application Layer

## A note on the use of these ppt slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) in substantially unaltered form, that you mention their source (after all, we'd like people to use our book!)
- If you post any slides in substantially unaltered form on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

All material copyright 1996-2009  
J.F Kurose and K.W. Ross, All Rights Reserved

# Chapter 2: Application layer

- r 2.1 Principles of network applications
- r 2.2 Web and HTTP
- r 2.3 FTP
- r 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- r 2.5 DNS
- r 2.6 P2P applications
- r 2.7 Socket programming with UDP
- r 2.8 Socket programming with TCP

# Web and HTTP

## First some jargon

- r Web page consists of objects
- r Object can be HTML file, JPEG image, Java applet, audio file,...
- r Web page consists of base HTML-file which includes several referenced objects
- r Each object is addressable by a URL
- r Example URL:

`www.someschool.edu/someDept/pic.gif`


host name

path name

# HTTP overview

## HTTP: hypertext transfer protocol

- r Web's application layer protocol
- r client/server model
  - ❖ *client*: browser that requests, receives, "displays" Web objects
  - ❖ *server*: Web server sends objects in response to requests



# HTTP overview (continued)

## Uses TCP:

- r client initiates TCP connection (creates socket) to server, port 80
- r server accepts TCP connection from client
- r HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- r TCP connection closed

## HTTP is "stateless"

- r server maintains no information about past client requests

# HTTP connections

## Nonpersistent HTTP

- r At most one object is sent over a TCP connection.

## Persistent HTTP

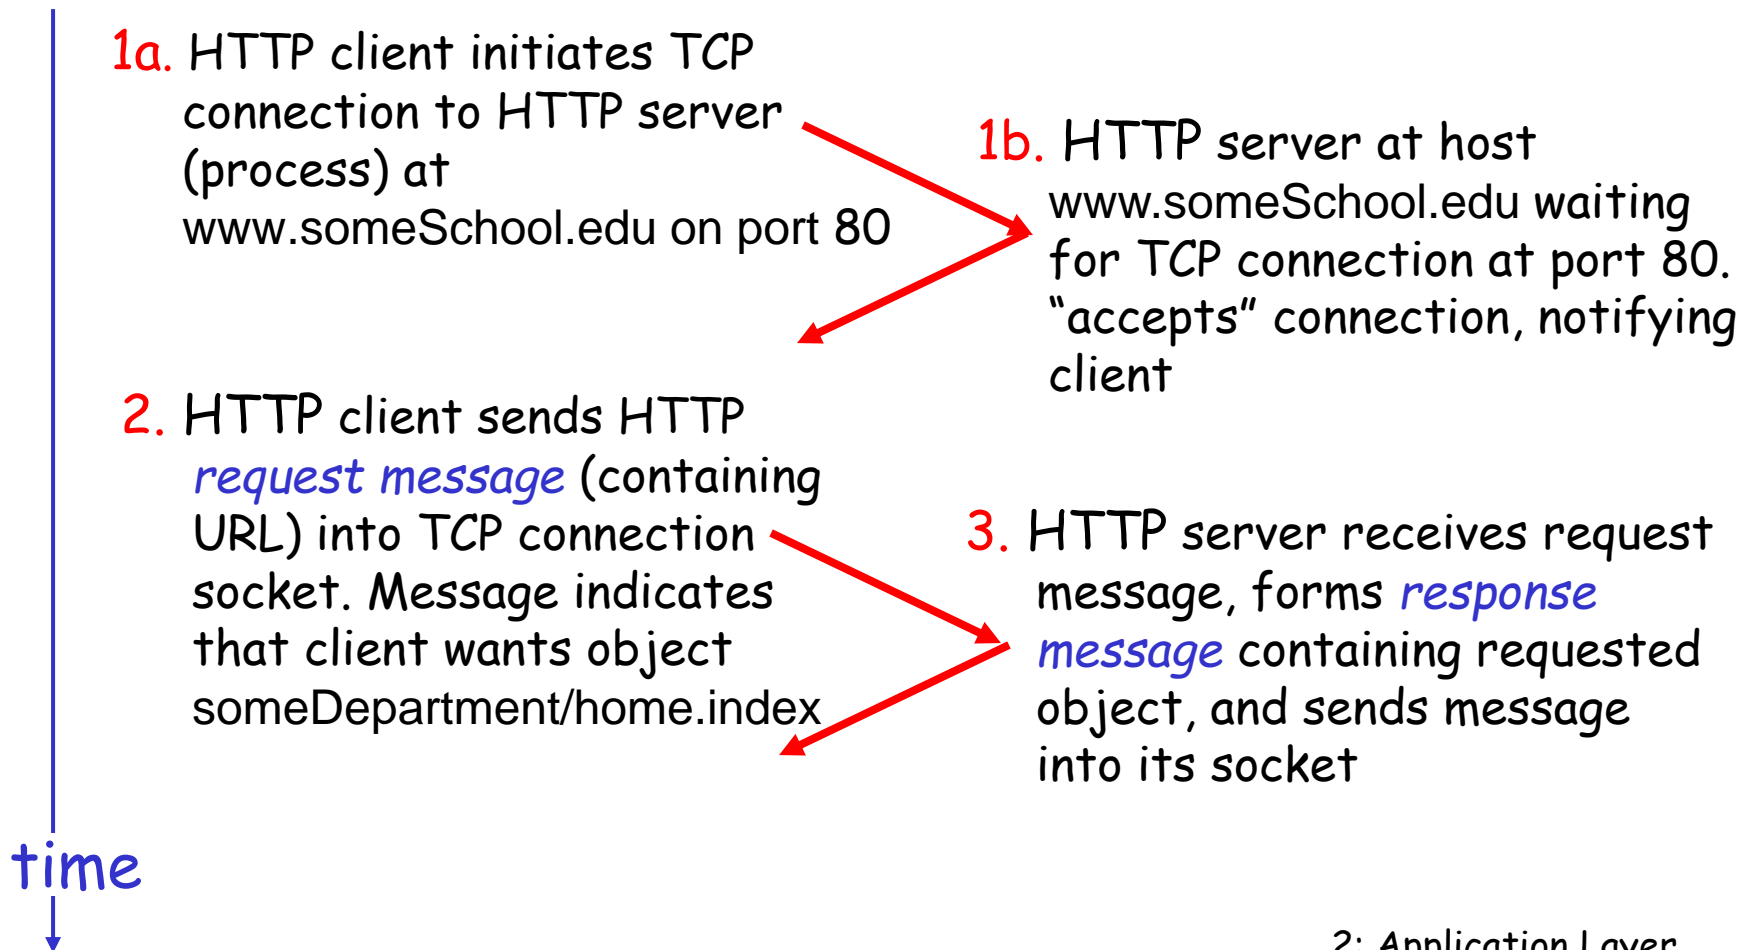
- r Multiple objects can be sent over single TCP connection between client and server.

# Nonpersistent HTTP

Suppose user enters URL

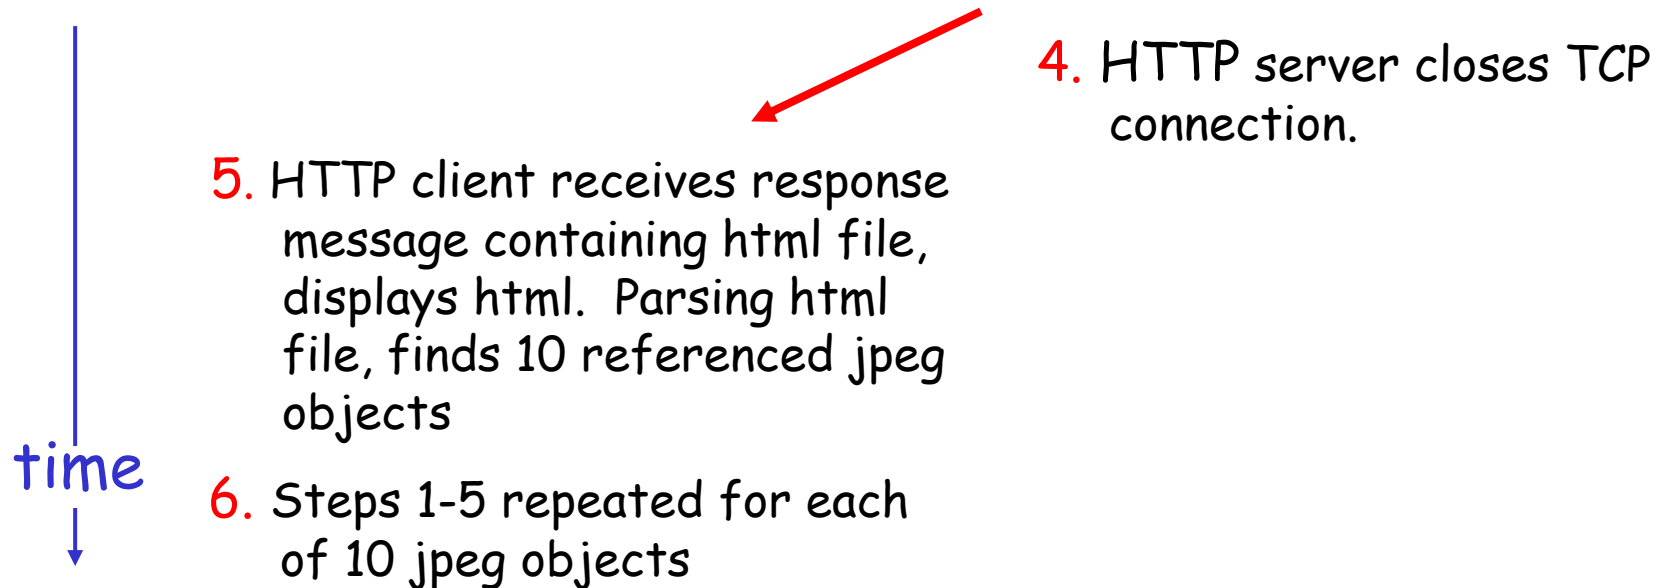
`www.someSchool.edu/someDepartment/home.index`

(contains text,  
references to 10  
jpeg images)





# Nonpersistent HTTP (cont.)




# Non-Persistent HTTP: Response time

**Definition of RTT:** time for a small packet to travel from client to server and back.

## Response time:

- r one RTT to initiate TCP connection
- r one RTT for HTTP request and first few bytes of HTTP response to return
- r file transmission time

**total =  $2RTT + \text{transmit time}$**



# Persistent HTTP

## Nonpersistent HTTP issues:

- r requires 2 RTTs per object
- r OS overhead for each TCP connection
- r browsers often open parallel TCP connections to fetch referenced objects

## Persistent HTTP

- r server leaves connection open after sending response
- r subsequent HTTP messages between same client/server sent over open connection
- r client sends requests as soon as it encounters a referenced object
- r as little as one RTT for all the referenced objects

# HTTP request message

- r two types of HTTP messages: *request, response*
- r **HTTP request message:**
  - ❖ ASCII (human-readable format)

request line  
(GET, POST,  
HEAD commands)


header  
lines

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
User-agent: Mozilla/4.0
Connection: close
Accept-language: fr
```

Carriage return,  
line feed  
indicates end  
of message

(extra carriage return, line feed)

# HTTP request message: general format



# Uploading form input

## Post method:

- r Web page often includes form input
- r Input is uploaded to server in entity body

## URL method:

- r Uses GET method
- r Input is uploaded in URL field of request line:

`www.somesite.com/animalsearch?monkeys&banana`

# Method types

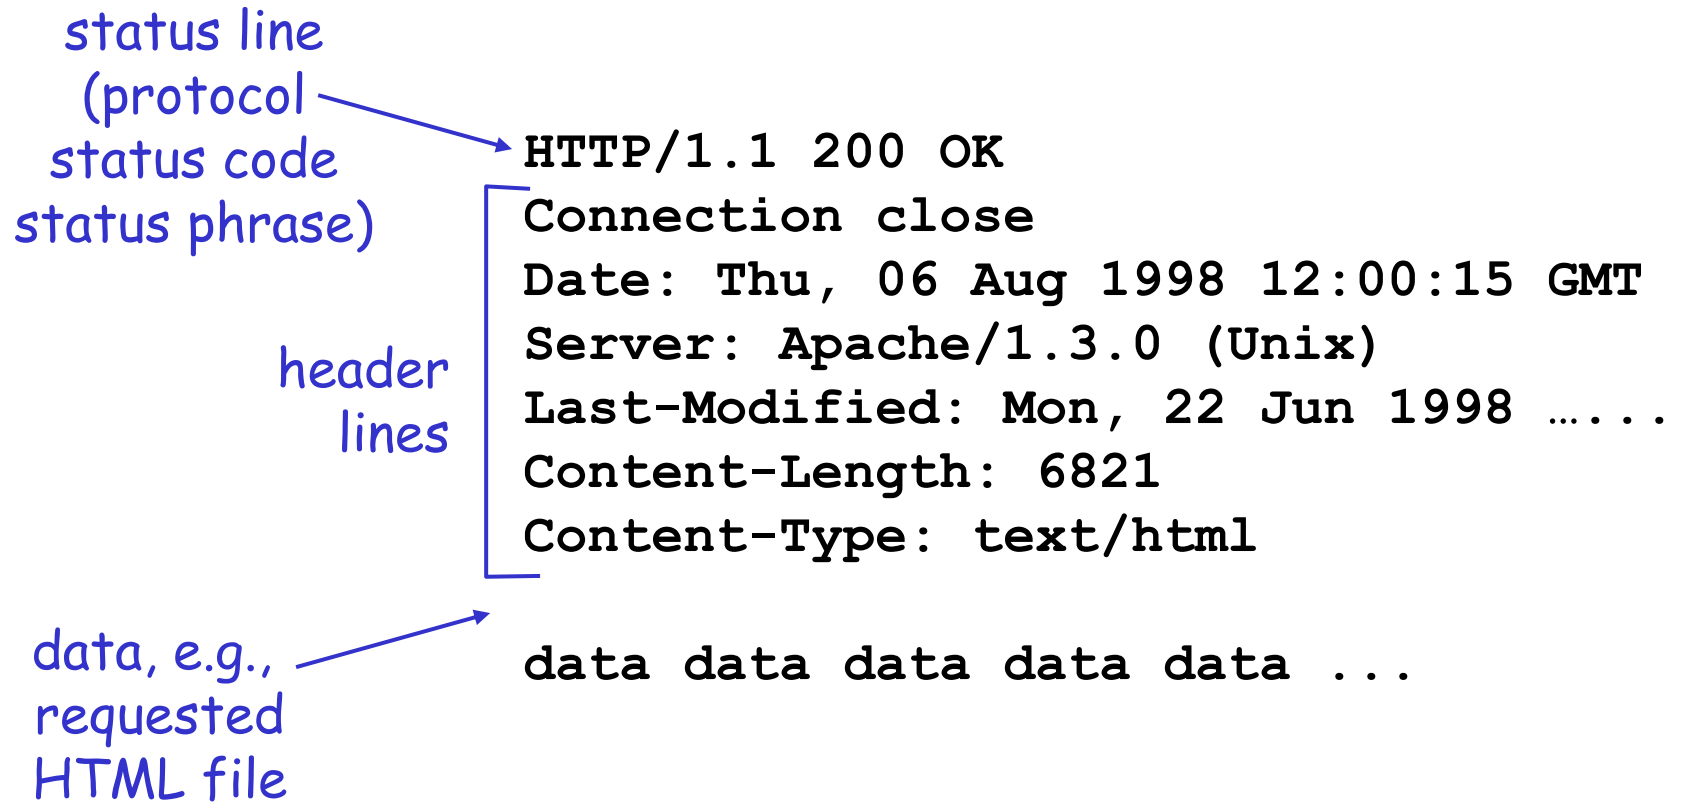
## HTTP/1.0

- r GET
- r POST

## HTTP/1.1

- r GET, POST
- r PUT
  - ❖ uploads file in entity body to path specified in URL field
- r DELETE
  - ❖ deletes file specified in the URL field

# HTTP response message





# HTTP response status codes

In first line in server->client response message.

A few sample codes:

## **200 OK**

- ❖ request succeeded, requested object later in this message

## **301 Moved Permanently**

- ❖ requested object moved, new location specified later in this message (Location:)

## **400 Bad Request**

- ❖ request message not understood by server

## **404 Not Found**

- ❖ requested document not found on this server

## **505 HTTP Version Not Supported**

# User-server state: cookies

Many major Web sites  
use cookies

## Four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

## Example:

- r Susan access Internet always from PC
- r visits specific e-commerce site for first time
- r when initial HTTP requests arrives at site, site creates:
  - ❖ unique ID
  - ❖ entry in backend database for ID

# Cookies: keeping "state" (cont.)

client

server



cookie file



usual http request msg

usual http response  
Set-cookie: 1678

usual http request msg  
cookie: 1678

usual http response msg

usual http request msg  
cookie: 1678

usual http response msg

Amazon server  
creates ID  
1678 for user


create  
entry

cookie-  
specific  
action

cookie-  
specific  
action

access

access



backend  
database

one week later:

# Cookies (continued)

## What cookies can bring:

- r authorization
- r shopping carts
- r recommendations
- r user session state  
(Web e-mail)

## How to keep "state":

- r protocol endpoints: maintain state at sender/receiver over multiple transactions
- r cookies: http messages carry state

— aside —


## Cookies and privacy:

- r cookies permit sites to learn a lot about you
- r you may supply name and e-mail to sites

# Web caches (proxy server)

**Goal:** satisfy client request without involving origin server

- r user sets browser:  
Web accesses via cache
- r browser sends all  
HTTP requests to  
cache
  - ❖ object in cache: cache  
returns object
  - ❖ else cache requests  
object from origin  
server, then returns  
object to client



# More about Web caching

- r cache acts as both client and server
- r typically cache is installed by ISP (university, company, residential ISP)

## Why Web caching?

- r reduce response time for client request
- r reduce traffic on an institution's access link.

# New uses of Web and HTTP

- r The Internet of Things
- r The Rest concept
  - ❖ The core IETF working group
  - ❖ RFC 6690

# Chapter 2: Application layer

- r 2.1 Principles of network applications
- r 2.2 Web and HTTP
- r 2.3 FTP
- r 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- r 2.5 DNS
- r 2.6 P2P applications
- r 2.7 Socket programming with UDP
- r 2.8 Socket programming with TCP



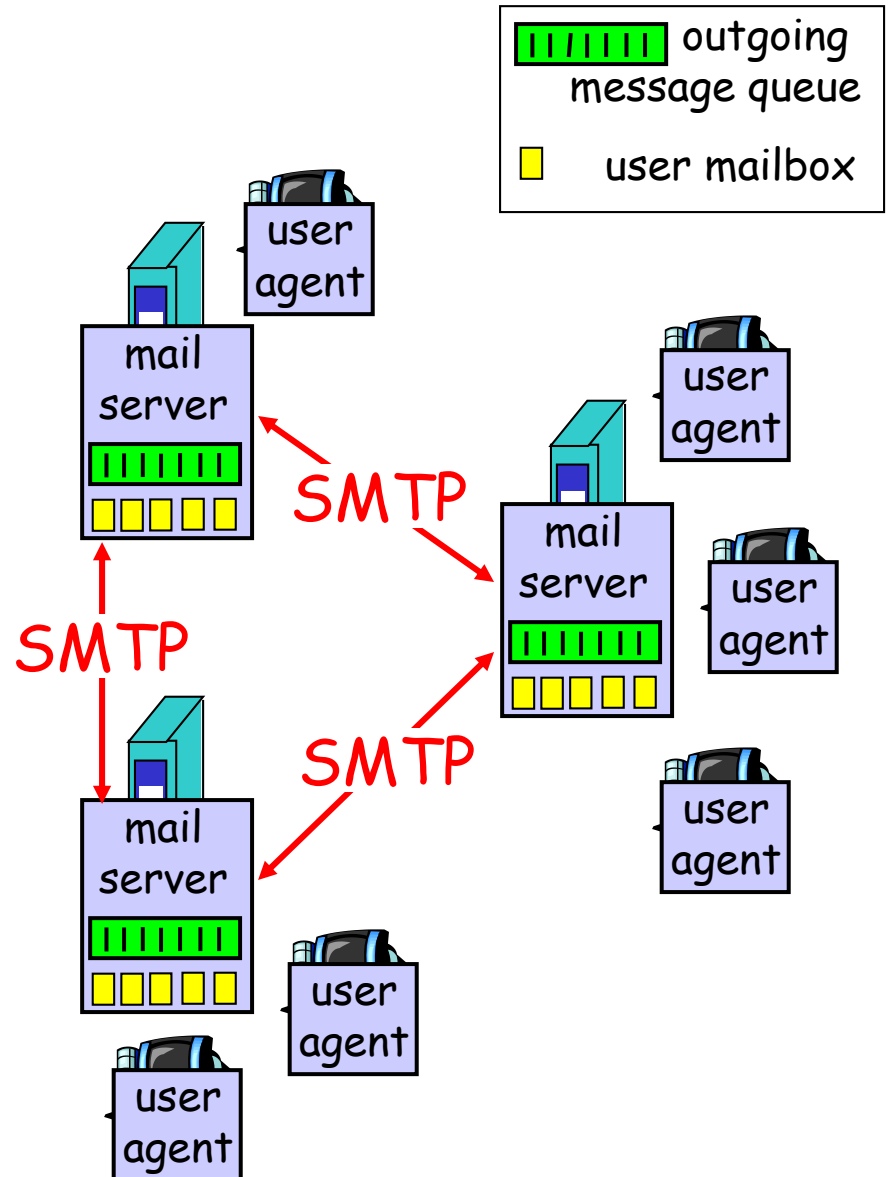
# Electronic Mail

## Three major components:

- r user agents
- r mail servers
- r simple mail transfer protocol: SMTP

## User Agent


- r a.k.a. "mail reader"
- r composing, editing, reading mail messages
- r e.g., Eudora, Outlook, elm, Mozilla Thunderbird
- r outgoing, incoming messages stored on server



# Electronic Mail: mail servers

## Mail Servers

- r **mailbox** contains incoming messages for user
- r **message queue** of outgoing (to be sent) mail messages
- r **SMTP protocol** between mail servers to send email messages
  - ❖ client: sending mail server
  - ❖ "server": receiving mail server

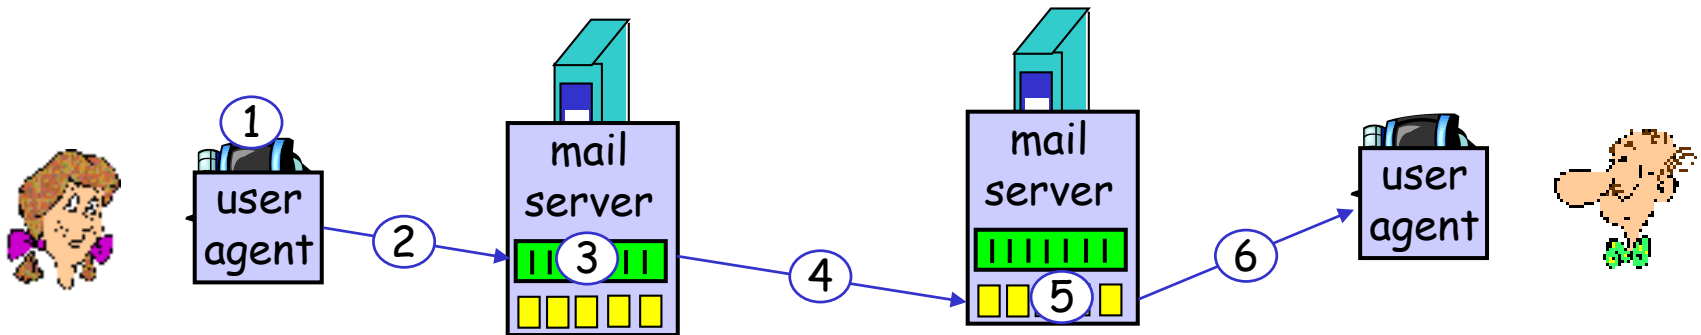


# Electronic Mail: SMTP [RFC 2821]

- r uses TCP to reliably transfer email message from client to server, port 25
- r direct transfer: sending server to receiving server
- r three phases of transfer
  - ❖ handshaking (greeting)
  - ❖ transfer of messages
  - ❖ closure
- r command/response interaction
  - ❖ **commands**: ASCII text
  - ❖ **response**: status code and phrase
- r messages must be in 7-bit ASCII

# Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message and "to"  
`bob@someschool.edu`
- 2) Alice's UA sends message to her mail server; message placed in message queue
- 3) Client side of SMTP opens TCP connection with Bob's mail server
- 4) SMTP client sends Alice's message over the TCP connection
- 5) Bob's mail server places the message in Bob's mailbox
- 6) Bob invokes his user agent to read message



# Sample SMTP interaction

```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

# SMTP: final words

- r SMTP uses persistent connections
- r SMTP requires message (header & body) to be in 7-bit ASCII
- r SMTP server uses CRLF.CRLF to determine end of message

## Comparison with HTTP:

- r HTTP: pull
- r SMTP: push
- r both have ASCII command/response interaction, status codes
- r HTTP: each object encapsulated in its own response msg
- r SMTP: multiple objects sent in multipart msg

# Mail message format

SMTP: protocol for exchanging email msgs

RFC 822: standard for text message format:


r header lines, e.g.,

- ❖ To:
- ❖ From:
- ❖ Subject:


*different from SMTP commands!*

r body

- ❖ the "message", ASCII characters only



# Mail access protocols



- r SMTP: delivery/storage to receiver's server
- r Mail access protocol: retrieval from server
  - ❖ POP: Post Office Protocol [RFC 1939]
    - authorization (agent <-->server) and download
  - ❖ IMAP: Internet Mail Access Protocol [RFC 1730]
    - more features (more complex)
    - manipulation of stored msgs on server
  - ❖ HTTP: gmail, Hotmail, Yahoo! Mail, etc.



# POP3 protocol

## authorization phase

r client commands:

- ❖ user: declare username
- ❖ pass: password

r server responses

- ❖ +OK
- ❖ -ERR

## transaction phase, client:

- r list: list message numbers
- r retr: retrieve message by number
- r dele: delete
- r quit

```
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

# POP3 (more) and IMAP

## More about POP3

- r Previous example uses "download and delete" mode.
- r Bob cannot re-read e-mail if he changes client
- r "Download-and-keep": copies of messages on different clients
- r POP3 is stateless across sessions

## IMAP

- r Keep all messages in one place: the server
- r Allows user to organize messages in folders
- r IMAP keeps user state across sessions:
  - ❖ names of folders and mappings between message IDs and folder name

# Chapter 2: Application layer

- r 2.1 Principles of network applications
- r 2.2 Web and HTTP
- r 2.3 FTP
- r 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- r 2.5 DNS
- r 2.6 P2P applications
- r 2.7 Socket programming with UDP
- r 2.8 Socket programming with TCP

# DNS: Domain Name System

**People:** many identifiers:

- ❖ SSN, name, passport #

**Internet hosts, routers:**

- ❖ IP address (32 bit) - used for addressing datagrams
- ❖ "name", e.g.,  
ww.yahoo.com - used by humans

**Q:** map between IP addresses and name ?

**Domain Name System:**

- r *distributed database*  
implemented in hierarchy of many *name servers*
- r *application-layer protocol*  
host, routers, name servers to communicate to *resolve* names (address/name translation)
  - ❖ note: core Internet function, implemented as application-layer protocol
  - ❖ complexity at network's "edge"

# DNS

## DNS services


- r hostname to IP address translation
- r host aliasing
  - ❖ Canonical, alias names
- r mail server aliasing
- r load distribution
  - ❖ replicated Web servers: set of IP addresses for one canonical name

## Why not centralize DNS?

- r single point of failure
- r traffic volume
- r distant centralized database
- r maintenance

*doesn't scale!*

# Distributed, Hierarchical Database

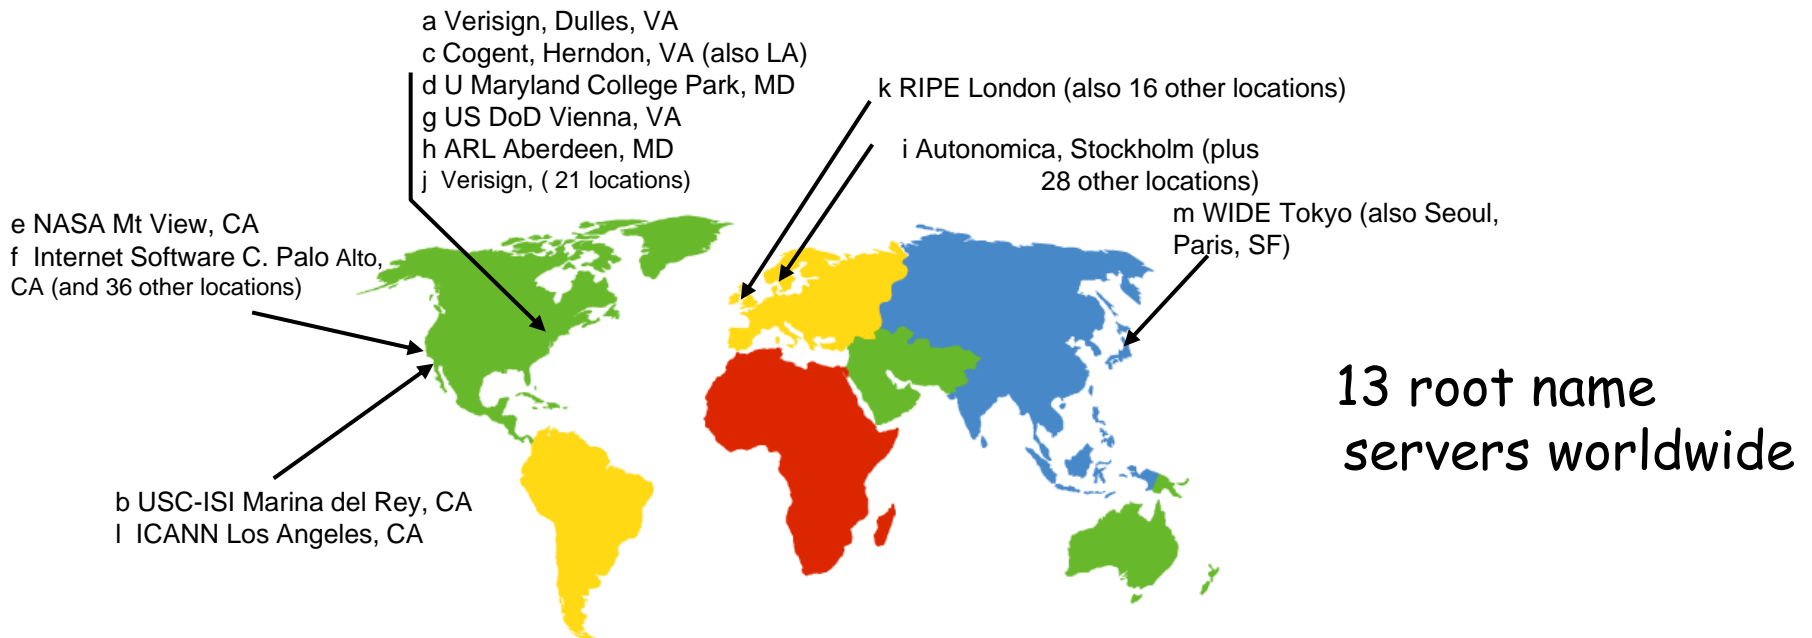


Client wants IP for www.amazon.com; 1<sup>st</sup> approx:

- r client queries a root server to find com DNS server
- r client queries com DNS server to get amazon.com DNS server
- r client queries amazon.com DNS server to get IP address for www.amazon.com

# DNS: Root name servers

- r contacted by local name server that can not resolve name
- r root name server:
  - ❖ contacts authoritative name server if name mapping not known
  - ❖ gets mapping
  - ❖ returns mapping to local name server



# TLD and Authoritative Servers

- r **Top-level domain (TLD) servers:**
  - ❖ responsible for com, org, net, edu, etc, and all top-level country domains uk, fr, ca, jp.
- r **Authoritative DNS servers:**
  - ❖ organization's DNS servers, providing authoritative hostname to IP mappings for organization's servers (e.g., Web, mail).
  - ❖ can be maintained by organization or service provider



# Local Name Server


- r does not strictly belong to hierarchy
- r each ISP (residential ISP, company, university) has one.
  - ❖ also called "default name server"
- r when host makes DNS query, query is sent to its local DNS server
  - ❖ acts as proxy, forwards query into hierarchy

# DNS name resolution example

- Host at cis.poly.edu wants IP address for gaia.cs.umass.edu

## iterated query:

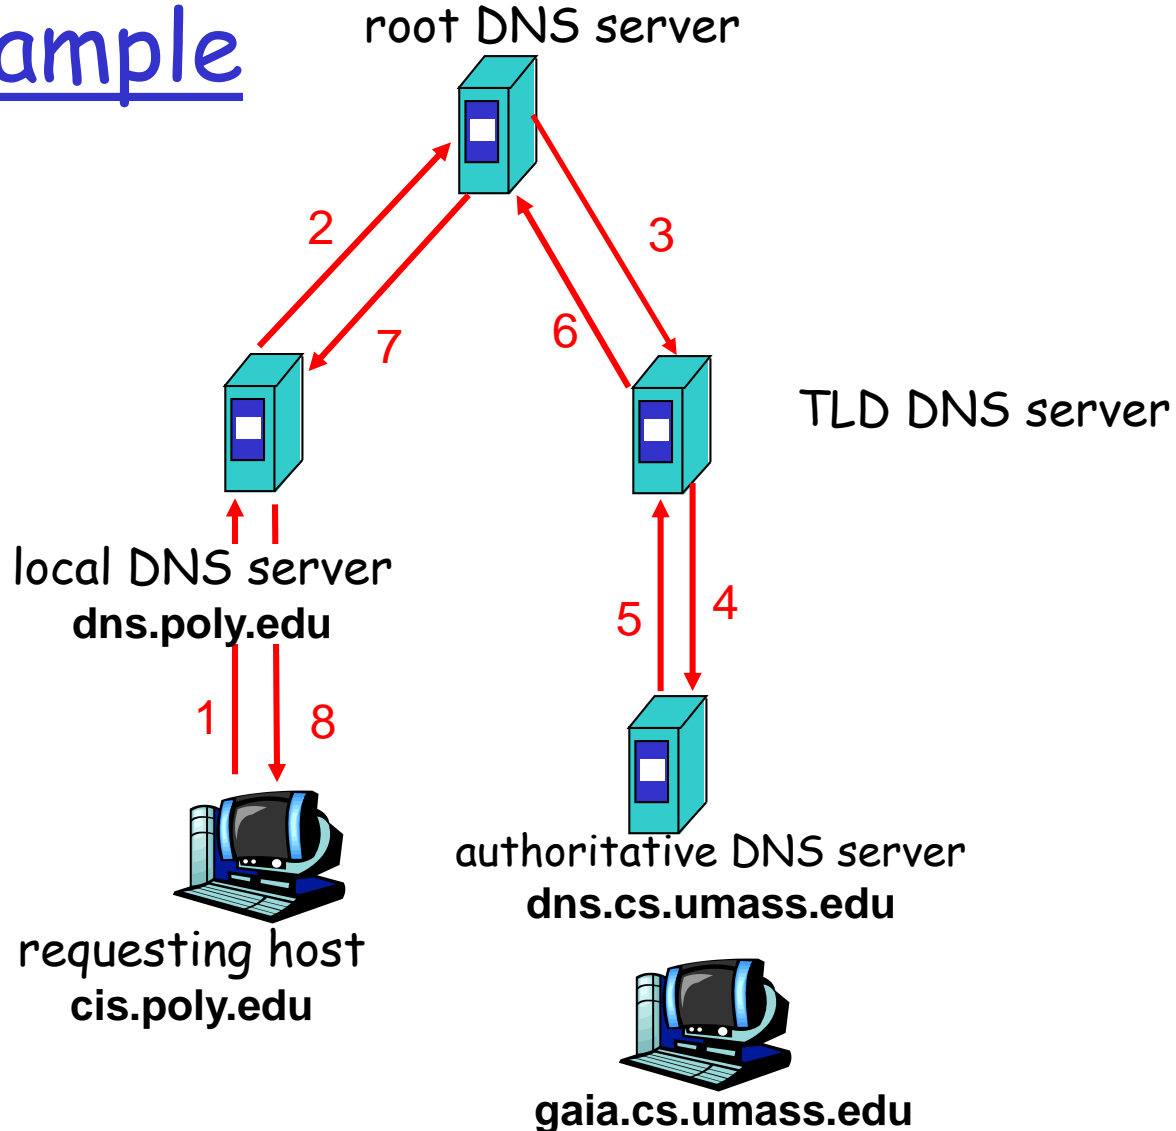
- contacted server replies with name of server to contact
- "I don't know this name, but ask this server"



# DNS name resolution example

## recursive query:

- r puts burden of name resolution on contacted name server
- r heavy load?



# DNS records

DNS: distributed db storing resource records (RR)

RR format: (name, value, type, ttl)

## r Type=A

- ❖ name is hostname
- ❖ value is IP address

## r Type=NS

- ❖ name is domain (e.g. foo.com)
- ❖ value is hostname of authoritative name server for this domain

## r Type=CNAME

- ❖ name is alias name for some "canonical" (the real) name  
www.ibm.com is really  
servereast.backup2.ibm.com
- ❖ value is canonical name

## r Type=MX

- ❖ value is name of mailserver associated with name

# DNS protocol, messages

DNS protocol : *query* and *reply* messages, both with same *message format*

## msg header

r **identification**: 16 bit #  
for query, reply to query  
uses same #

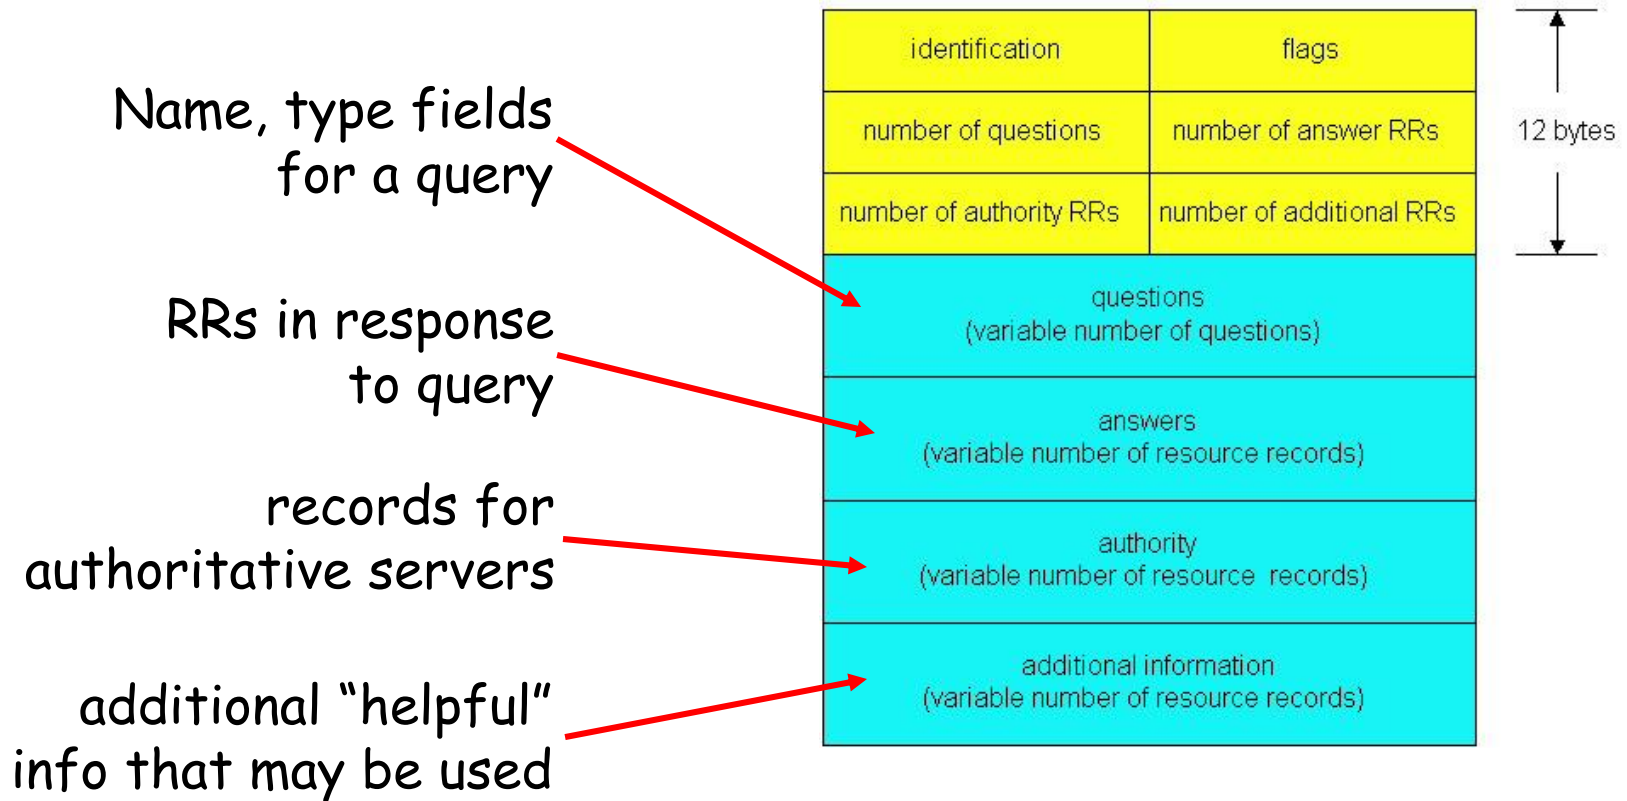
r **flags**:

- ❖ query or reply
- ❖ recursion desired
- ❖ recursion available
- ❖ reply is authoritative

identification	flags
number of questions	number of answer RRs
number of authority RRs	number of additional RRs
questions (variable number of questions)	
answers (variable number of resource records)	
authority (variable number of resource records)	
additional information (variable number of resource records)	

↑  
12 bytes  
↓

# DNS protocol, messages



# Inserting records into DNS

- r example: new startup "Network Utopia"
- r register name networkutopia.com at DNS *registrar* (e.g., Network Solutions)
  - ❖ provide names, IP addresses of authoritative name server (primary and secondary)
  - ❖ registrar inserts two RRs into com TLD server:

```
(networkutopia.com, dns1.networkutopia.com, NS)
(dns1.networkutopia.com, 212.212.212.1, A)
```

- r create authoritative server Type A record for www.networkutopia.com; Type MX record for networkutopia.com
- r *How do people get IP address of your Web site?*